

目录

1. Priority queues
 1. Inserting and delmaxing
 2. A Priority Queue ADT
 3. Client: PQ sort
 4. Three implementations of the priority queue ADT
 1. PQ ADT implementation using an unordered array
 2. PQ ADT implementation using an ordered array
 1. Performance: ordered vs un-ordered
 3. PQ implementation using a heap
 1. Performance: unordered vs ordered vs heap-based
 2. Heap Sort

Priority queues

In a queue:

- items are processed in the order of arrival
- you could say that the first item in the queue always has highest priority

In a priority queue:

- items are assigned *priorities*: they are ordered by priority
- the item with the highest priority (value) is processed first
- the priority may have nothing to do with the order of arrival

A priority queue is a generalization of a stack and a queue

- can implement a stack using a priority queue
- can implement a queue using a priority queue

Like queues and stacks, *priority queues* can be implemented as ADTs and implemented by:

- arrays or
- linked lists

A priority queue supports 2 abstract operations:

- **insert**: a new item (equivalent to *push* for a stack, and *qush* for a queue)
- **delmax**: delete the item with the largest key (equivalent *pop*)

You should also be able to

- create a priority queue
- is the priority queue empty

Inserting and delmaxing

Consider the input sequence $E X A \# M \# \# \#$

- where a letter defines an insert operation (push)
 - if implemented as a function, it puts the element in the data structure and returns a *void*
- a # defines a delete operation (pop) and return the value
 - if implemented as a function, it deletes and an element in the data structure and returns the element

If the ADT was a stack, then # deletes the *top* element:

<u>operation</u>	<u>return</u>	<u>stack</u>
<i>insert E</i>	void	E
<i>insert X</i>	void	EX
<i>insert A</i>	void	EXA
<i>pop</i>	A	EX
<i>insert M</i>	void	EXM
<i>pop</i>	M	EX
<i>pop</i>	X	E
<i>pop</i>	E	~

If the ADT was a queue, then # deletes the *bottom* element:

<u>operation</u>	<u>return</u>	<u>queue</u>
<i>insert E</i>	void	E
<i>insert X</i>	void	EX
<i>insert A</i>	void	EXA
<i>pop</i>	E	XA
<i>insert M</i>	void	XAM
<i>pop</i>	X	AM
<i>pop</i>	A	M
<i>pop</i>	M	~

If the ADT is a priority_queue then # deletes the *maximum* (or *minimum* element)

- we refer to it as **delmax** (or **delmin**)
- the sequence of operations is:

<u>operation</u>	<u>return</u>	<u>priority_queue</u>
<i>insert E</i>	void	E
<i>insert X</i>	void	EX
<i>insert A</i>	void	EXA
<i>delmax</i>	X	EA
<i>insert M</i>	void	EAM
<i>delmax</i>	M	EA
<i>delmax</i>	E	A
<i>delmax</i>	A	~

Notice:

- in the ADT we do not now in what order elements are stored in the abstract data structure
- *delmax* gets the maximum element, but we do not know how *delmax* finds this element

Consider another example: *P R I O # R # # I # T # Y # # # Q U E # # # U # E*

- The order that the elements will be deleted is:
 - first hash deletes the first *R*
 - second hash deletes the second *R*
 - third hash deletes *P*
 - fourth hash deletes *O*
 - etc

A Priority Queue ADT

Typically an ADT allows the user to:

- **create** an abstract data structure
- **insert** an element into the structure
 - stacks and queues use *push* and *qush*
- **delete** an element from the structure
 - stacks and queues use *pop*
- check whether the structure **is empty**

A generic interface could be:

切换行号显示

```
1 // pq.h: ADT interface for a priority queue
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 typedef struct pqRep *PQ;
6
7 PQ  createPQ(int size);
8 void insertPQ(PQ q, int it);
9 int  delMaxPQ(PQ q);
10 int  isEmptyPQ(PQ q);
```

Other possible operations in the interface could be:

切换行号显示

```
1 void  updatePQ(PQ q, int i, int j); // change element value from i
to j
2 void  deletePQ(PQ q);              // remove the PQ completely
3 PQ    joinPQ(PQ q1, PQ q2);       // concatenate 2 PQs
4
```

Client: PQ sort

Here is a client of the ADT, called **pqSort()**, that sorts an array of integers.

切换行号显示

```

1  /* pqSort.c: use a priority queue to sort an array of integers
2      into descending order
3  */
4  #include "pq.h"
5
6  int main() {
7      int a[] = {41, 2, 58, 156, 360, 81, 260, 74, 167, 13};
8      int length = sizeof(a)/sizeof(a[0]);
9
10     PQ q = createPQ(length);
11     printf("Array: ");
12     for (int i = 0; i < length; i++) {
13         printf("%d ", a[i]);
14         insertPQ(q, a[i]);
15     }
16     printf("\nSorted: ");
17     while (!isEmptyPQ(q)) {
18         printf("%d ", delMaxPQ(q));
19     }
20     putchar('\n');
21     return EXIT_SUCCESS;
22 }

```

This program:

- inserts all the items into the priority queue
- delMaxs items from the priority queue in descending order

If you want the minimum element, rewrite *delmaxPQ()* into *delMinPQ()*

Is the program fast?

- that depends on how *insertPQ()* and *delMaxPQ()* are implemented
- *PQ* is an ADT, so the client does not know

Is the program efficient in space used?

- *No*, because
 - each item in the array is inserted into the priority queue
 - ... which is also an array, so 2 arrays are being used!
- it would be better if *pqSort()* was *in-place*

Three implementations of the priority queue ADT

We can implement the interface functions using different concrete data structures:

1. an unordered array
 - items are simply appended to the array as they are entered
2. an ordered array
 - items are inserted into their correct position in an array
3. a heapified array (i.e. a heap)
 - items are added to the array that is subsequently heapified

Note, all implementations use arrays, but they differ in how/where they store the items in the array.

PQ ADT implementation using an unordered array

切换行号显示

```

1 // pqUA.c: priority queue implementation for pq.h using an
  unordered array
2 #include "pq.h"
3
4 struct pqRep {
5     int nItems; // actual count of Items
6     int *items; // array of Items
7     int maxsize; // maximum size of array
8 };
9
10 PQ createPQ(int size) {
11     PQ q = malloc(sizeof(struct pqRep)); // make room for the
  structure
12     if (q == NULL) {
13         fprintf(stderr, "out of memory\n");
14         exit(0);
15     }
16
17     q->items = malloc(size * sizeof(int)); // make room for the
  array
18     if (q->items == NULL) {
19         fprintf(stderr, "out of memory\n");
20         exit(0);
21     }
22     q->nItems = 0; // we have no items yet
23     q->maxsize = size; // remember the maxsize
24     return q; // return the initial
  PQ
25 }
26
27 void insertPQ(PQ q, int it) {
28     if (q == NULL) {
29         fprintf(stderr, "priority queue not initialised\n");
30         exit(1);
31     }
32     if (q->nItems == q->maxsize) {
33         fprintf(stderr, "priority queue full\n");
34         exit(1);
35     }
36     q->items[q->nItems] = it; // UNORDERED ARRAY, so put item at
  the end
37     q->nItems++; // increment the 'counter'
38 }
39
40 int delMaxPQ(PQ q) { // UNORDERED, so need to linear search for
  max item
41     if (q == NULL) {
42         fprintf(stderr, "delmaxPQ: priority queue not
  initialised\n");
43         exit(1);
44     }
45     if (q->nItems == 0) {
46         fprintf(stderr, "priority queue empty\n");
47         exit(1);
48     }
49     int *array = q->items;
50     int last = q->nItems-1; // items occupy places 0 ..
  last
51     int max = 0; // assume initially item at
  max=0 has largest key
52     for (int i = 1; i <= last; i++){
53         if (array[max] < array[i]){ // now compare with every other
  item
54             max = i; // whenever we find a better
  one, update max
55         }
56     }

```

```

57     int retval = array[max];           // save the max item
58     array[max] = array[last];          // overwrite max location with
last item
59     q->nItems--;                        // decrease the number of items
60     return retval;                     // return the max element
61 }
62
63 int isEmptyPQ(PQ q) {
64     int empty = 0;
65     if (q == NULL) {
66         fprintf(stderr, "isEmptyPQ: priority queue not
initialised\n");
67     }
68     else {
69         empty = q->nItems == 0;
70     }
71     return empty;
72 }

```

Compile and run the client *pqSort()* with this ADT:

```

prompt$ gcc pqUA.c pqSort.c
prompt$ ./a.out
Array: 41 2 58 156 360 81 260 74 167 13
Sorted: 360 260 167 156 81 74 58 41 13 2

```

Notice the descending order, produced by calling *delMaxPQ()* in the client.

PQ ADT implementation using an ordered array

We could have implemented the priority queue using an ADT based on an Ordered Array.

- this will make it much easier to find the maximum element ...
- ... but when we insert, we will need to do more work

切换行号显示

```

1 // pqOA.c: priority queue implementation for pq.h using an ordered
array
2 #include "pq.h"
3
4 struct pqRep {
5     int nItems; // actual count of Items
6     int *items; // array of Items
7     int maxsize; // maximum size of array
8 };
9
10 PQ createPQ(int size) {
11     PQ q = malloc(sizeof(struct pqRep)); // make room for the
structure
12     if (q == NULL) {
13         fprintf(stderr, "out of memory\n");
14         exit(0);
15     }
16     q->items = malloc(size * sizeof(int)); // make room for the
array
17     if (q->items == NULL) {
18         fprintf(stderr, "out of memory\n");
19         exit(0);
20     }
21     q->nItems = 0; // we have no items yet
22     q->maxsize = size; // remember the maxsize
23     return q; // return the initial
PQ
24 }

```

```

25
26 void insertPQ(PQ q, int it) {
27     if (q == NULL) {
28         fprintf(stderr, "priority queue not initialised\n");
29         exit(1);
30     }
31     if (q->nItems == q->maxsize) {
32         fprintf(stderr, "priority queue full\n");
33         exit(1);
34     }
35     int *array = q->items;
36     int last = q->nItems;
37     int i;
38     for (i=0; i<last && array[i]<it; i++) {
39         ; // find location of item == it
40     }
41     int j;
42     for (j = last; j>i; j--){ // starting at last and go down to
i
43         array[j] = array[j-1]; // shift items up
44     }
45     array[i] = it; // now insert item 'it' at i
46     q->nItems++; // increase the count
47 }
48
49 int delMaxPQ(PQ q) {
50     if (q == NULL) {
51         fprintf(stderr, "priority queue not initialised\n");
52         exit(1);
53     }
54     if (q->nItems == 0) {
55         fprintf(stderr, "priority queue empty\n");
56         exit(1);
57     }
58     q->nItems--;
59     return q->items[q->nItems];
60 }
61
62 int isEmptyPQ(PQ q) {
63     int empty = 0;
64     if (q == NULL) {
65         fprintf(stderr, "isEmptyPQ: priority queue not
initialised\n");
66     }
67     else {
68         empty = q->nItems == 0;
69     }
70     return empty;
71 }

```

Compile and run the client *pqSort.c* with this ADT:

```

prompt$ gcc pqOA.c pqSort.c
prompt$ ./a.out
Array: 41 2 58 156 360 81 260 74 167 13
Sorted: 360 260 167 156 81 74 58 41 13 2

```

- the output is the same as for the unordered array

Example of use of the unordered and ordered array *PQ* implementations:

<u>Operation</u>	<u>Ordered</u>	<u>Unordered</u>	returned value
insertPQ E	E	E	-
insertPQ X	EX	EX	-

insertPQ A	AEX	EXA	-
insertPQ M	AEMX	EXAM	-
delMaxPQ	AEM	EMA	X

Note:

- you cannot see from 'outside the ADT' which order the items are in
- all you can see is the results of the *delMaxPQ* operation that removes them from the *PQ*
 - above, the same item 'X' is returned in both cases

Performance: ordered vs un-ordered

The performance of the 2 priority queue ADT implementations can be compared:

<u>Implementation</u>	<u>insertPQ</u>	<u>delMaxPQ</u>
ordered array	$O(N)$	$O(1)$
unordered array	$O(1)$	$O(N)$

Hence, we have poor performance either with an *insertPQ* or a *delMaxPQ*

- we know that, using a heap, we can *insert* and *delMax* in $O(\log(n))$
- **can we use a heap to implement a priority queue ADT?**

PQ implementation using a heap

A heap is also an array, like ordered and unordered

- (actually it could be some other data structure, but in this course it will always be an array)
- it contains the same elements, but in a different order than ordered and unordered
- the difference is the effect of *HOP* and *CTP*

切换行号显示

```

1 // pqHP.c: priority queue implementation for pq.h using a heap
2 #include "pq.h"
3
4 // 'static' means these functions are for local use only
5 static void fixDown(int *, int, int);
6 static void fixUp(int *, int);
7
8 // Priority queue implementation using an unordered array
9 struct pqRep {
10     int nItems; // actual count of Items
11     int *items; // array of Items
12     int maxsize; // maximum size of array
13 };
14
15 PQ createPQ(int size) {
16     PQ q = malloc(sizeof(struct pqRep)); // make room for the
structure
17     if (q == NULL) {
18         fprintf(stderr, "out of memory\n");
19         exit(0);
20     }
21     q->items = malloc((size+1) * sizeof(int)); // make room for the
array
22     if (q->items == NULL) { // size+1 because heap

```



```

1..size
23     fprintf(stderr, "out of memory\n");
24     exit(0);
25 }
26 q->nItems = 0; // we have no items yet
27 q->maxsize = size; // remember the maxsize
28 return q; // return the initial
PQ
29 }
30
31 void insertPQ(PQ q, int it) {
32     if (q == NULL) {
33         fprintf(stderr, "priority queue not initialised\n");
34         exit(1);
35     }
36     if (q->nItems == q->maxsize) {
37         fprintf(stderr, "priority queue full\n");
38         exit(1);
39     }
40     q->nItems++; // adding another item
41     q->items[q->nItems] = it; // put the item at the end
42     fixUp(q->items, q->nItems); // fixUp all the way to the
root
43     return;
44 }
45
46 int delMaxPQ(PQ q) {
47     if (q == NULL) {
48         fprintf(stderr, "priority queue not initialised\n");
49         exit(1);
50     }
51     if (q->nItems == 0) {
52         fprintf(stderr, "priority queue empty\n");
53         exit(1);
54     }
55     int retval = q->items[1]; // this is the item we want
to return
56     q->items[1] = q->items[q->nItems]; // overwrite root by last
item
57     q->nItems--; // we are decreasing heap
size by 1
58     fixDown(q->items, 1, q->nItems); // fixDown the new root
59     return retval;
60 }
61
62 int isEmptyPQ(PQ q) {
63     int empty = 0;
64     if (q == NULL) {
65         fprintf(stderr, "isEmptyPQ: priority queue not
initialised\n");
66     }
67     else {
68         empty = q->nItems == 0;
69     }
70     return empty;
71 }
72
73 // fix up the heap for the 'new' element child
74 void fixUp(int *heap, int child) {
75     while (child > 1 && heap[child/2] < heap[child]) {
76         int swap = heap[child]; // if parent < child, do a
swap
77         heap[child] = heap[child/2];
78         heap[child/2] = swap;
79         child = child/2; // become the parent
80     }
81     return;
82 }
83

```

```

84 // force value at a[par] into correct position
85 void fixDown(int *heap, int par, int len) {
86     int finished = 0;
87     while (2*par<=len && !finished) { // as long as you are within
bounds
88         int child = 2*par;           // the first child is here
89         if (child<len && heap[child]<heap[child+1]) {
90             child++;                  // choose larger of two children
91         }
92         if (heap[par]<heap[child]) { // if node is smaller than this
child ...
93             int swap = heap[child]; // if parent < child, do a swap
94             heap[child] = heap[child/2];
95             heap[child/2] = swap;
96             par = child;              // ... and become this child
97         }
98         else {
99             finished = 1;            // else we do not have to go any
further
100         }
101     }
102     return;
103 }

```

Compile and run:

```

prompt$ gcc pqHP.c pqSort.c
prompt$ ./a.out
Array: 41 2 58 156 360 81 260 74 167 13
Sorted: 360 260 167 156 81 74 58 41 13 2

```

- this is the same output as for the unordered and ordered arrays

Performance: unordered vs ordered vs heap-based

Cost of operations in a heap:

- tree height is $\log(n)$
- each insert/delete requires at most $\log(n)$ compares/swaps on a path from root to leaf
- complexity is $O(\log(n))$ for *insertPQ()* and *delMaxPQ()*

For just one operation:

<u>Implementation</u>	<u>insertPQ</u>	<u>delMaxPQ</u>
ordered array	$O(n)$	$O(1)$
unordered array	$O(1)$	$O(n)$
heap array	$O(\log(n))$	$O(\log(n))$

Multiply by n for an array of length n .

Heap Sort

The client *pqSort()* using the heap array implementation is heap sort

- ... although it can be implemented more efficiently

The client:

- stores the unsorted list as a heap (ADT)
- repeatedly deletes the maximum element (ADT)
- heapifying each time (ADT)
- Heap sort is **not adaptive** (*this is bad*)
 - means that it is not faster if the data is already partially ordered
 - even if data is ordered, *fixUp* and *fixDown* may still need to go from root to leaf
- Heap sort is **not stable** (*this is bad*)
 - means that terms with the same value may change order
- Heap sort is **in-place** ... (this is good)
 - ... but *pqSort()* is not

Heap sort is $O(n \log(n))$ (best, average and worst case behaviour)

- merge sort, quick sort as well, but quick sort has worst case $O(n^2)$ (rarely happens)
- Heap sort is not recursive (*this is good*), merge sort and quick sort are recursive
- heapsort is faster than mergesort (by a constant factor) and more space efficient
 - ... although in a distributed environment (lots of machines) merge sort is better
- quick sort is generally faster than heap sort (by a constant factor) for most data sets

If speed is critical, and adaptivity & stability do not matter, heap sort is usually preferred.

 Heap sort animation

PQs (2019-07-11 16:59:40由AlbertNymeyer编辑)