

目录

1. Trees and Binary Search Trees
 1. Types of trees
 2. Binary Trees
 1. Height of a binary tree
 2. Depth of a binary tree
2. Binary Search Tree
 1. 1. Data structure for a binary tree
 2. Comparing Binary Search Trees and Heaps
 3. Searching in BSTs
 4. Creating a node in a BST
 5. Freeing a node in a BST
 6. Inserting a node in a BST
 1. Example: putting the basic tree operations together
 7. Printing a BST
3. More functions on a BST
 1. Count the number of nodes in a BST
 2. Find the height of a BST
 3. How balanced is a BST?
 1. Example: Putting it all together
4. Deleting a node from a BST
 1. 1. node is a leaf
 2. node has 1 child
 3. node has 2 children
 2. Delete a node code
5. BST operations complexity analysis
 1. Summary of operations

Trees and Binary Search Trees

(Chapter 5.4 - 5.7 Sedgewick)

Trees are data structures, like arrays and linked lists.

Trees are like doubly-linked lists: nodes contain data and multiple links to other nodes.

- Nodes are:
 - internal, and have links to other nodes, called their children
 - external, called *leaves* or *terminals*, and have no links to other nodes
- Every node has a parent node, except one, which is called the *root* node
- The *descendants* of a node consist of all the nodes reachable on a path from that node.
- Children with the same parents are called *siblings*

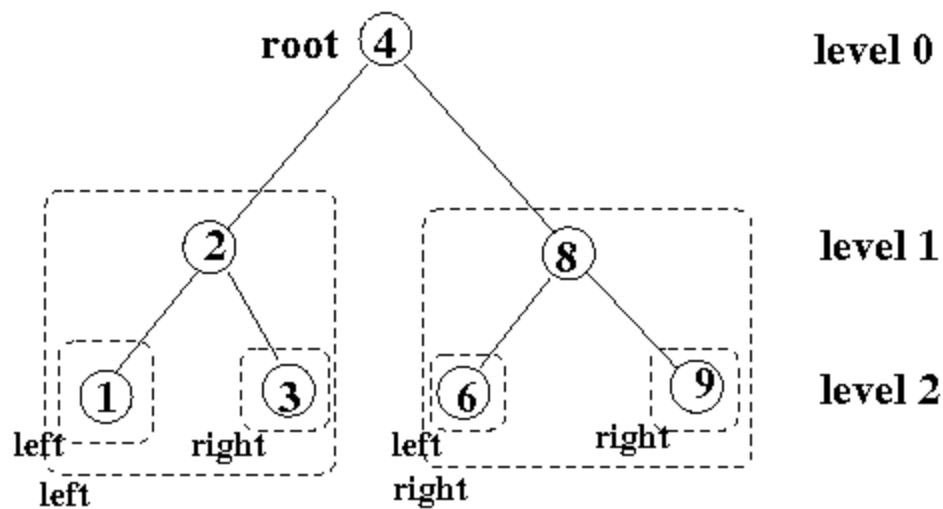
Formally, a tree is an acyclic graph in which each child has at most one parent.

Connections between nodes are called edges or links

- A **path** is a set of connected edges
- We usually consider paths that go only one way i.e. only up or down
- **Height** is the length of the longest path from the root
 - so the root node is at height 0
 - children of the root node are at height 1
- **Node level** or **depth** is the path length from the root to the node

- Depth of the root is 0
- Hierarchy of trees and subtrees
 - assuming 2 children for each internal node:
 - tree at left child is called a *left subtree*
 - tree at right child is called a *right subtree*

Binary tree



Types of trees

Assume we have a tree with internal nodes and leaves, and where each node has a data value

- *ordered tree*
 - children are in order (left and right children have order)
- *binary tree*
 - each internal node has at most 2 children
- *full binary tree*
 - each internal node has exactly 2 children
- *perfect binary tree*
 - binary tree in which all leaves are at the same depth
- *ordered binary tree*
 - left subtree values \leq parent value
 - right subtree values \geq parent value
 - great for searching: e.g.
 - search for smallest: keep going left down the tree
 - search for largest: keep going right down the tree
 - search for a specific element: use 'classic' binary search
 - also called a ***binary search tree***
- *full m-ary tree*
 - each internal node has exactly m children


Binary Trees

Binary trees can be

- **balanced**
 - tree has minimal height for the given number of nodes
 - number of nodes in the left and right subtrees differ by at most 1

- **degenerate**

- tree with maximal height (i.e. every parent has 1 child)

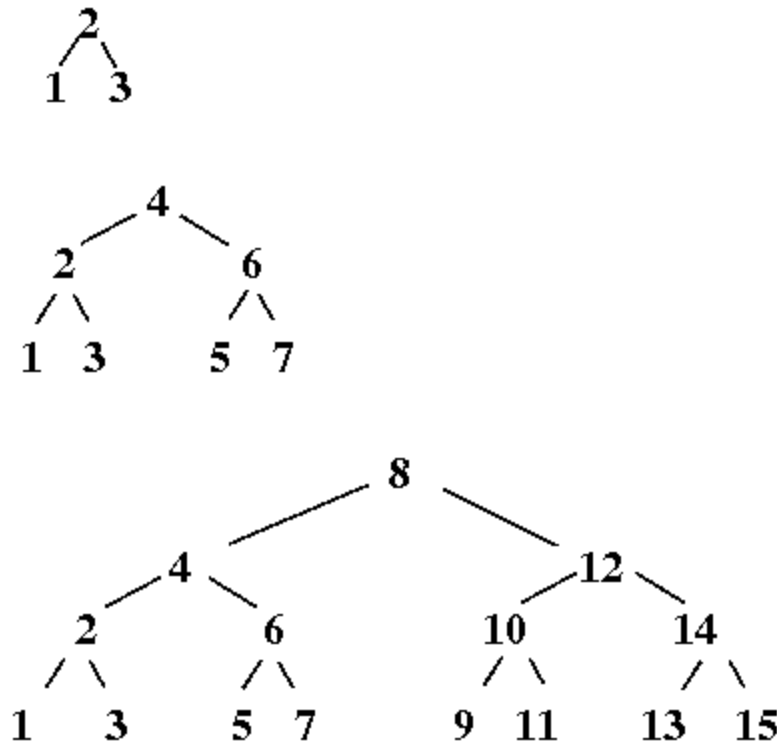
 Wikipedia: binary trees

Height of a binary tree

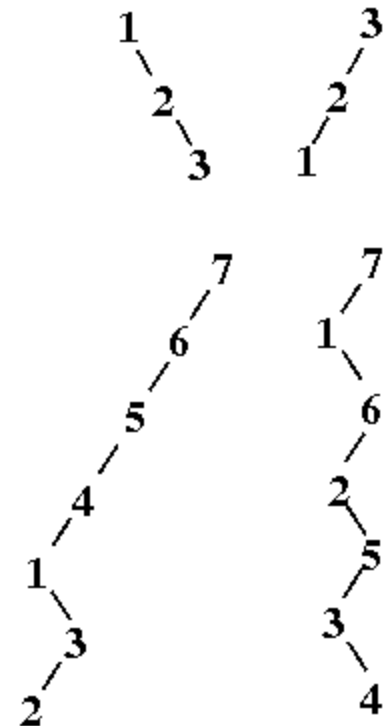
What is the height of a binary tree consisting of n nodes:

- what was the definition of height?
 - the length of the longest path
- what is the maximum height?
 - the tree is degenerate
 - height is $n-1$
- what is the minimum height?
 - the tree is balanced
 - height is $\lg(n)$

balanced



degenerate



Number	Balanced Height	Degenerate Height
3	1	2
7	2	6
15	3	14
n	$\lg n$	$n-1$

Depth of a binary tree

The depth of a node x is the length (in edges) of the path from x to the root. Computationally,

- if n is a root node then $depth(n) = 0$
- else $depth(n) = 1 + depth(parent_of\ n)$

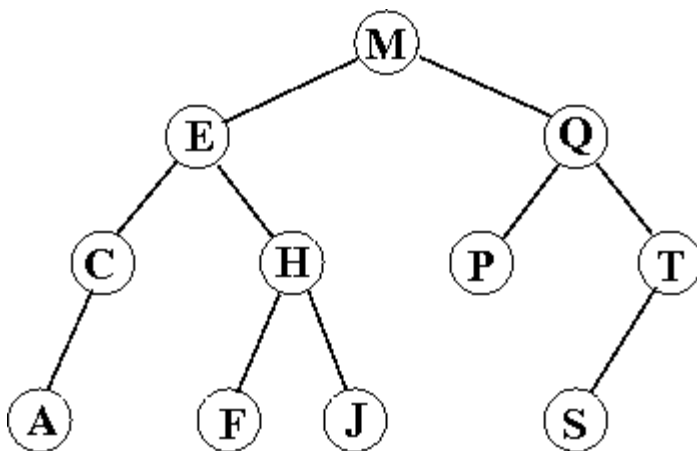
The maximum depth of any node in a tree is the height of the tree

Binary Search Tree

A BST is a tree where for every ('parent') node:

- * if the node has a left child, its key is smaller than the key of the node
- * if it has a right child, its key is larger than or equal to the key of the node

Example of a BST:



- notice they are ordered from left to right if you 'abstract away' the height (i.e. flatten the tree)

Data structure for a binary tree

切换行号显示

```

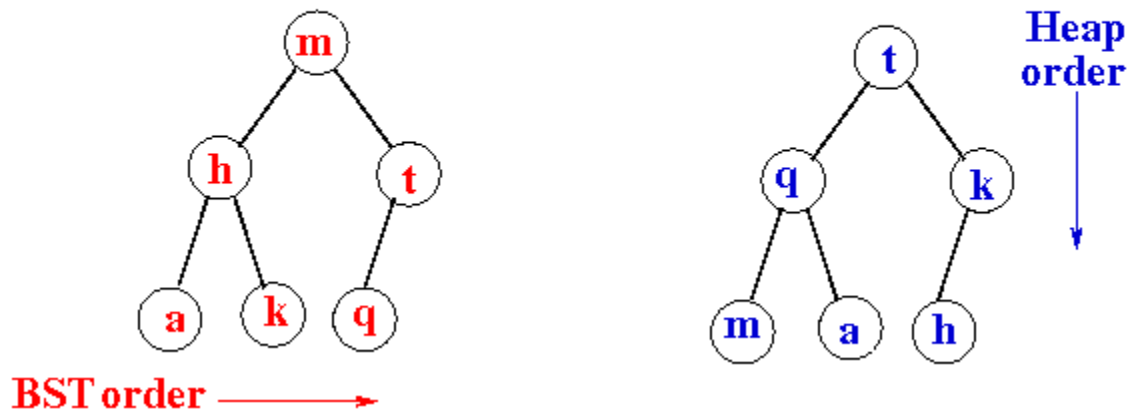
1  typedef struct node *Tree;
2  struct node {
3      int data;
4      Tree left;
5      Tree right;
6  };
  
```

Comparing Binary Search Trees and Heaps

- **Heaps** are trees with *top-to-bottom ordering*
 - satisfy the *Complete Tree Property*
- **Binary Search Trees** are trees with *left-to-right ordering*
 - there is NO *Complete Tree Property* for BSTs
 - ... they can be degenerate!
 - ... hence cannot be implemented as arrays. We must use linked lists.

Here is an example of a BST and a heap for the same input:

Insert order: m t h q a k



A BST satisfies the property:

- for node with key k
 - all node keys in left subtree $< k$
 - all node keys in right subtree $\geq k$
 - property applies to all nodes in the BST

BSTs can be great for searching

- if n is the size of the input, and the height of the BST is $\ln(n)$ then
 - binary search performs as $O(\ln(n))$

... but the bad news is that BSTs can be degenerate

- the BST then has height n
- this is the worst-case behaviour
- binary search then performs as $O(n)$
 - this is just linear search, which is much slower (*remember the Sydney phone book analogy!*)

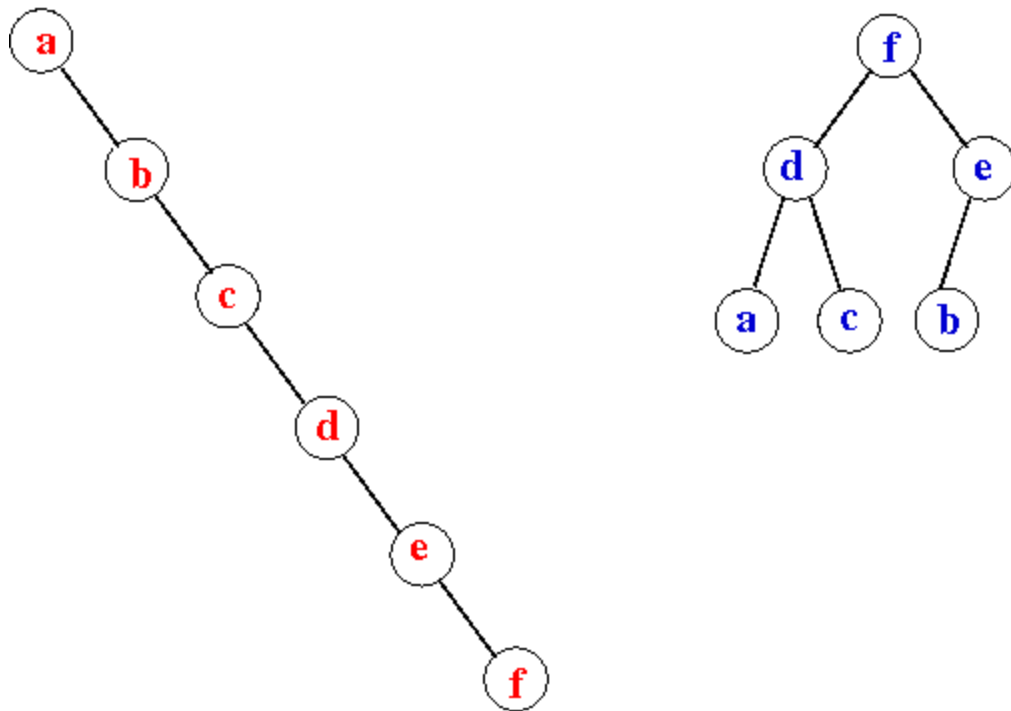
We construct a BST as we read the elements

- we cannot control the order of the input ...
- in fact, we will build a degenerate tree if the input is ordered

On the left is the result if we simply insert the nodes as we read

- on the right is what we would like the result to be

Insert order: a b c d e f



Searching in BSTs

A BST is a perfect data structure to do binary search

Reminder: what is a *binary search*?

- Prerequisite:
 - the items in a sequence must be sorted
- It is a *divide and conquer* technique
 - split the data into 2 parts
 - determine to which part the item belongs
 - recurse down until arrive at the base case
 - which is either NULL (element not found)
 - or the element itself

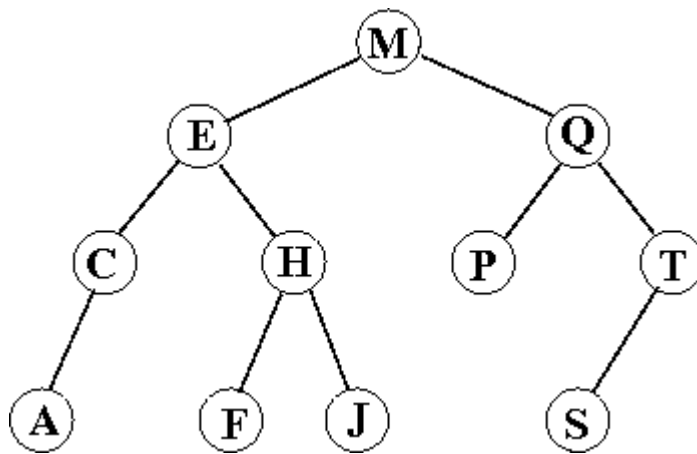
Assume a node contains a single key (and maybe some more data).

- a key is just the data that we want to order by, and search for

Basic idea:

If the value of the item is less than the item in the current node, then go left, otherwise go right

For example, given the following BST:



To search for item F

- $searchTree(M, F)$
 - go left and $searchTree(E, F)$
 - go right and $searchTree(H, F)$
 - go left and $searchTree(F, F)$
 - *success*

If the search was for item G , then we would have had the sequence

- $searchTree(M, G)$
 - go left and $searchTree(E, G)$
 - go right and $searchTree(H, G)$
 - go left and $searchTree(F, G)$
 - F is a leaf and $F \neq G$ so *failure*

切换行号显示

```

1 int searchTree(Tree t, int v){ // Search for the node with value v
2     int ret = 0;
3     if (t != NULL) {
4         if (v < t->data) {
5             ret = searchTree(t->left, v);
6         }
7         else if (v > t->data) {
8             ret = searchTree(t->right, v);
9         }
10        else { // v == t->data
11            ret = 1;
12        }
13    }
14    return ret;
15 } // returns non-zero if found, zero otherwise
16

```

Creating a node in a BST

Just like a linked list, we must:

- call *malloc()* to create the tree node
- initialise the data
- initialise the pointers

切换行号显示

```

1 typedef struct node *Tree;
2 struct node {
3     int data;
4     Tree left;
5     Tree right;
6 };
7
8 Tree createTree(int v) {
9     Tree t;
10    t = malloc(sizeof(struct node));
11    if (t == NULL) {
12        fprintf(stderr, "Out of memory\n");
13        exit(1);
14    }
15    t->data = v;
16    t->left = NULL;
17    t->right = NULL;
18    return t;
19 }

```

Freeing a node in a BST

The pointers in a BST node point to other nodes

- we need to follow the pointers to the last node, and work backwards freeing nodes
- otherwise we will have (severe) memory leaks

In the following code, we recurse down the tree, and call *free(t)* for each node from the bottom up

切换行号显示

```

1 void freeTree(Tree t) { // free in postfix fashion
2     if (t != NULL) {
3         freeTree(t->left);
4         freeTree(t->right);
5         free(t);
6     }
7     return;
8 }

```

Inserting a node in a BST

Trees seem to be linked lists with 2 links instead of 1(?)

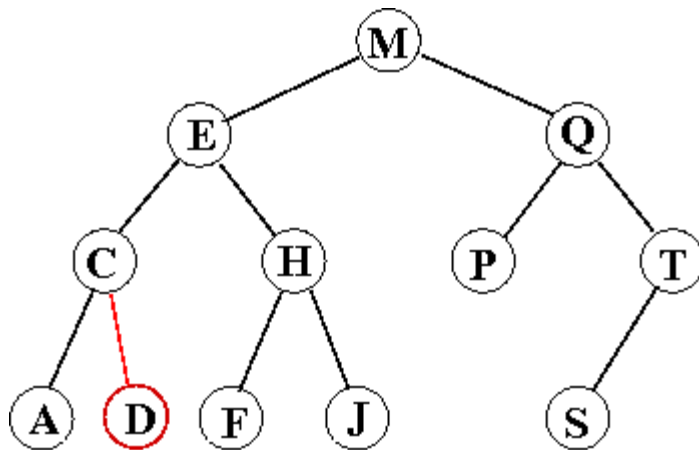
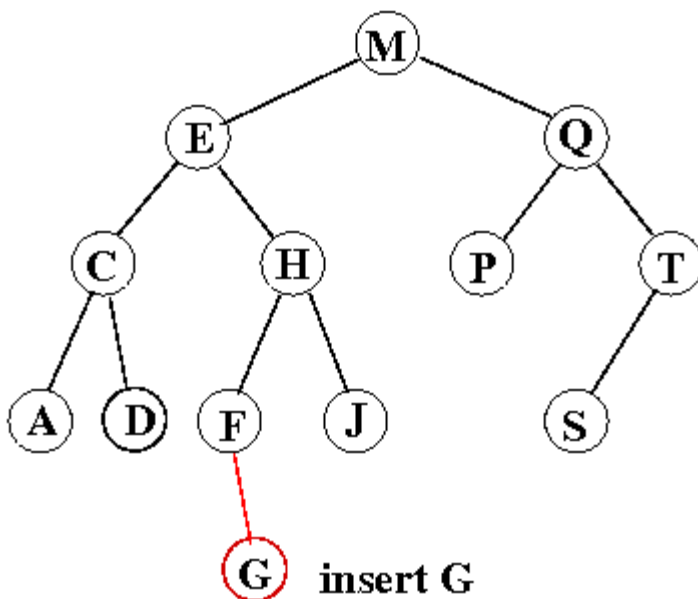
- **traversing** is similar (just follow the left or right link)
- **insertion?**
 - obvious in a linked list, but what strategy is used in a BST?
- **deletion?**
 - obvious in a linked list, but what happens to the 'children' in a BST?

In a BST, when we insert a new node:

- it **always becomes a leaf** (impossible to insert a non-leaf node)
- it **always maintains the ordering** of the tree
 - it must be on the left of all nodes larger than it
 - it must be on the right of all nodes smaller than or equal to it

Algorithm:

- i. follow the path from the root towards the leaves as though we were searching for the item
- ii. when we get to a NULL, we have found our insertion point
 - the NULL must be either the left or right child of a node
 - if the value is smaller than the current node, then the NULL is the left node
 - otherwise it becomes the right child (in our implementation, duplicates go on the right)
- iii. create a node for the item, and link it in by replacing the NULL with it

**insert D****insert G**

Finding the insertion point can be done recursively:

切换行号显示

```

1 Tree insertTree(Tree t, int v) {
2     if (t == NULL) {
3         t = createTree(v);
4     }
5     else {
6         if (v < t->data) {
7             t->left = insertTree (t->left, v);
8         }
9         else {
10            t->right = insertTree (t->right, v);
11        }

```

```

12     }
13     return t;
14 }

```

or we can do it iteratively:

切换行号显示

```

1 Tree insertTreeI(Tree t, int v) { // An iterative version of the
above
2     if (t == NULL) {
3         t = createTree(v);
4     }
5     else { // t != NULL
6         Tree parent = NULL; // remember the parent to link in new
child
7         Tree step = t;
8         while (step != NULL) { // this is the iteration
9             parent = step;
10            if (v < step->data) {
11                step = step->left;
12            }
13            else {
14                step = step->right;
15            }
16        } // step == NULL
17        if (v < parent->data) {
18            parent->left = createTree(v);
19        }
20        else {
21            parent->right = createTree(v);
22        }
23    }
24    return t;
25 }

```

The order of the input can make a huge difference in the structure of the BST

For example consider the input values 1, 2, 3 and 4

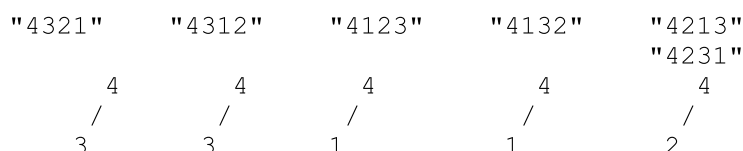
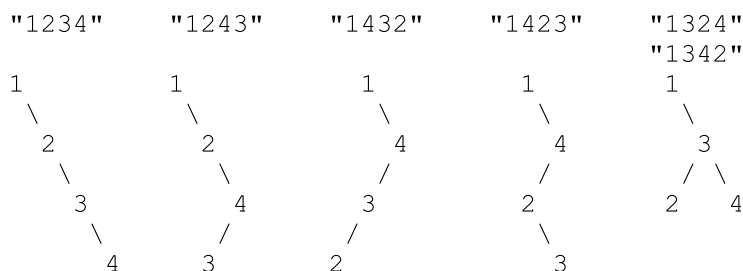
- There are $4 \times 3 \times 2$ possible orders of these 4 numbers:

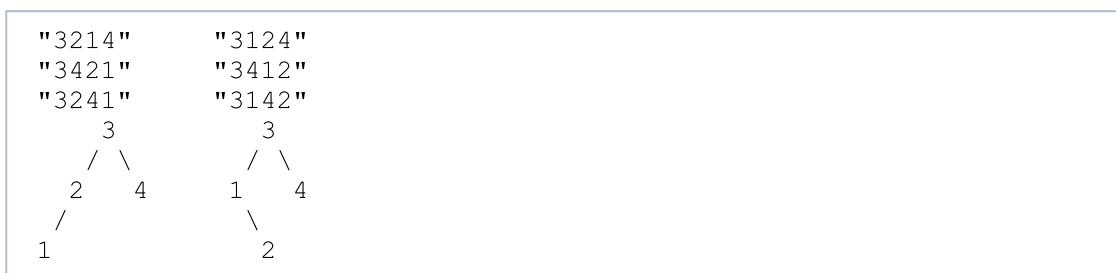
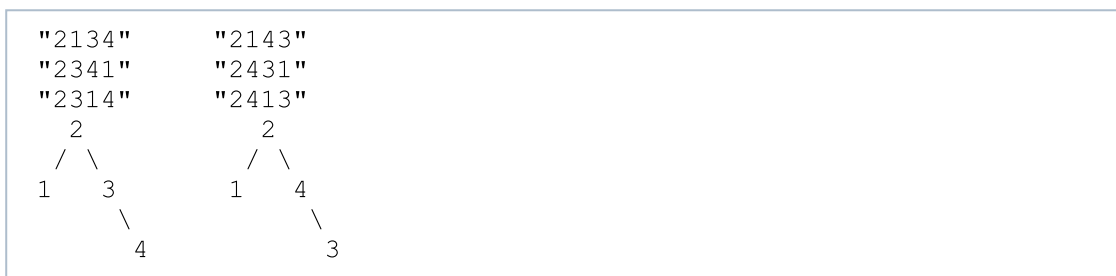
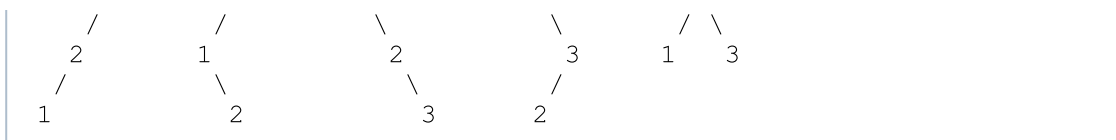
```

1234 1243 1324 1342 1423 1432
2134 2143 2314 2341 2413 2431
3124 3142 3214 3241 3412 3421
4123 4132 4213 4231 4312 4321

```

What is the BST that results from each of these 24 'orders':





Notice, different input can result in the same BST.

Example: putting the basic tree operations together

切换行号显示

```

1 // basic.c: insert nodes into a BST, print the tree and free all
nodes
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 typedef struct node *Tree;
6 struct node {
7     int data;
8     Tree left;
9     Tree right;
10 };
11
12 Tree insertTree (Tree, int);
13 Tree createTree (int);
14 void printTree (Tree);
15 void freeTree (Tree);
16
17 int main(void) {
18     Tree t;
19
20     t = createTree (7);
21     t = insertTree(t, 8);
22     t = insertTree(t, 6);
23     t = insertTree(t, 5);
24     t = insertTree(t, 4);
25     t = insertTree(t, 3);
26     t = insertTree(t, 2);
27     t = insertTree(t, 1);
28     printTree(t);
29     putchar('\n');
30     freeTree(t);
31     return EXIT_SUCCESS;
32 }
33
34 Tree insertTree(Tree t, int v) {
35     if (t == NULL) {
36         t = createTree(v);

```

```

37     }
38     else {
39         if (v < t->data) {
40             t->left = insertTree (t->left, v);
41         }
42         else {
43             t->right = insertTree (t->right, v);
44         }
45     }
46     return t;
47 }
48
49 Tree createTree (int v) {
50     Tree t = NULL;
51
52     t = malloc (sizeof(struct node));
53     if (t == NULL) {
54         fprintf(stderr, "Memory is exhausted: exiting\n");
55         exit(1);
56     }
57     t->data = v;
58     t->left = NULL;
59     t->right = NULL;
60     return t;
61 }
62
63 void printTree(Tree t) { // not the final version
64     if (t != NULL) {
65         printTree (t->left);
66         printf ("%d ", t->data);
67         printTree (t->right);
68     }
69     return;
70 }
71
72 void freeTree(Tree t) { // free in postfix fashion
73     if (t != NULL) {
74         freeTree(t->left);
75         freeTree(t->right);
76         free(t);
77     }
78     return;
79 }

```

There is no input to the program: the values are hard-coded into the program

- insert 7, 8, 6, 5, 4, 3, 2, 1

The output is:

```
1 2 3 4 5 6 7 8
```

Notice that 1 was the last input value, but 1 is the first output

- ... because it is the most left-most descendent of the root

Let's consider how we can draw a tree as a 2D structure (and not a sequence of numbers)

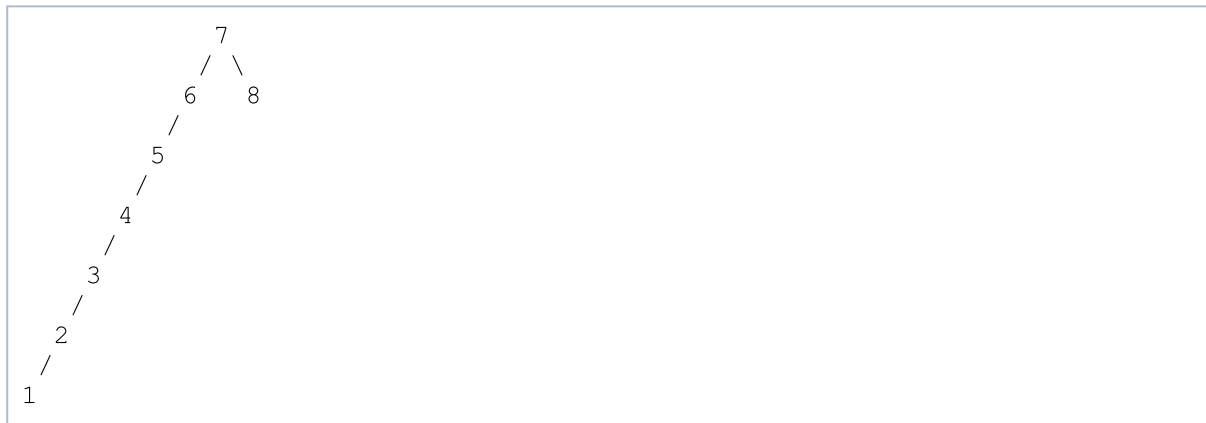
Printing a BST

The *printTree()* function actually converts the 2-D tree into 1-D **infix notation**:

- left child -- parent -- right child

- the order is correct, reading left to right, but there is no structure

What does the 'real' BST above look like?



Our *infix printTree()* above recursively descends the tree as follows:

```

left child of 7 is 6
  left child of 6 is 5
    left child of 5 is 4
      left child of 4 is 3
        left child of 3 is 2
          left child of 2 is 1
            left child of 1 is NULL
              parent; print '1'          <== 1
                right child of 1 is NULL
                  parent; print '2'      <== 2
                    right child is NULL
                      parent; print '3'  <== 3
                        right child is NULL
                          parent; print '4' <== 4
                            right child is NULL
                              parent; print '5' <== 5
                                right child is NULL
                                  parent; print '6' <== 6
                                    right child is NULL
                                      parent; print '7' <== 7
                                        right child of 7 is 8
                                          left child of 8 is NULL
                                            parent; print '8'          <== 8
                                              right child of 8 is NULL
  
```

Can we draw the BST properly (as a 2-D structure) based on the above?

What about printing the tree 'on its side'?



- notice the order is 1...8
- each new 'generation' goes on a new line
- the deeper we go, the bigger the indent
 - pass a 'depth' parameter down the tree to print leading spaces or tabs

Here is the result:

切换行号显示

```

1 void printTree(Tree t, int depth) { // extra depth parameter
2     if (t != NULL) {
3         depth++;
4         printTree (t->left, depth);
5         for (int i=1; i<depth; i++){ // 'depth'*whitespace
6             putchar('\t');
7         }
8         printf ("%d\n", t->data); // node to print
9         printTree (t->right, depth);
10    }
11    return;
12 }
```

The result is that we indent every node by n tabs, where n is the depth of the node

- '1' is at level 6
- '2' is at level 5
- ...
- '6' and '8' are at level 1
- '7' is at level 0

This generates the BST lying on its side shown above.

More functions on a BST

Count the number of nodes in a BST

切换行号显示

```

1 int count(Tree t){
2     int counttree = 0;
3     if (t != NULL) {
4         counttree = 1 + count(t->left) + count(t->right);
5     }
6     return counttree;
7 }
```

Find the height of a BST

First define a helper function to find the maximum of two numbers:

切换行号显示

```

1 int max(int a, int b){
2     if (a >= b) {
3         return a;
4     }
5     return b;
6 }
```

Now a function that returns the height of a BST:

切换行号显示

```

1  int height(Tree t){
2      int heighttree = -1;
3      if (t != NULL){
4          heighttree = 1 + max(height(t->left), height(t->right));
5      }
6      return heighttree;
7  }

```

How balanced is a BST?

How do the number of nodes in the left sub-tree and the right sub-tree compare?

切换行号显示

```

1  int balance (Tree t){ // calculates the difference between left
and right
2      int diff = 0;
3
4      if (t != NULL) {
5          diff = count(t->left) - count(t->right); // count
declared elsewhere
6          if (diff < 0) {
7              diff = -diff;
8          }
9      }
10     return diff;
11 }

```

Example: Putting it all together

- create a tree and output its balance, height and count, and free the data structure

切换行号显示

```

1  // unbalanced.c: create and check the balance of a tree
2  #include <stdio.h>
3  #include <stdlib.h>
4
5  typedef struct node *Tree;
6  struct node {
7      int data;
8      Tree left;
9      Tree right;
10 };
11
12 void printTree(Tree, int); // print a BST with indentation
13 Tree createTree(int); // create a BST with root 'v'
14 Tree insertTree(Tree, int); // insert a node 'v' into a BST
15 void freeTree(Tree); // give the memory back to the heap
16
17 int count(Tree);
18 int balance(Tree);
19 int height(Tree);
20
21 int main(void) {
22     Tree t;
23
24     t = createTree(7);
25     t = insertTree(t, 8);
26     t = insertTree(t, 6);
27     t = insertTree(t, 5);
28     t = insertTree(t, 4);
29     t = insertTree(t, 3);

```

```
30     t = insertTree(t, 2);
31     t = insertTree(t, 1);
32     printTree (t, 0);
33     printf("Balance = %d\n", balance(t));
34     printf("Height = %d\n", height(t));
35     printf("Count = %d\n", count(t));
36
37     freeTree(t);
38     return EXIT_SUCCESS;
39 }
40 void printTree(Tree t, int depth) {
41     if (t != NULL) {
42         depth++;
43         printTree (t->left, depth);
44         int i;
45         for (i=1; i<depth; i++){
46             putchar('\t');
47         }
48         printf ("%d\n", t->data);
49         printTree (t->right, depth);
50     }
51     return;
52 }
53
54 Tree createTree (int v) {
55     Tree t;
56     t = malloc(sizeof(struct node));
57     if (t == NULL) {
58         fprintf(stderr, "Out of memory\n");
59         exit(1);
60     }
61     t->data = v;
62     t->left = NULL;
63     t->right = NULL;
64     return t;
65 }
66
67 Tree insertTree(Tree t, int v) {
68     if (t == NULL) {
69         t = createTree(v);
70     }
71     else {
72         if (v < t->data) {
73             t->left = insertTree (t->left, v);
74         }
75         else {
76             t->right = insertTree (t->right, v);
77         }
78     }
79     return t;
80 }
81
82 int count(Tree t){
83     int counttree = 0;
84     if (t != NULL) {
85         counttree = 1 + count(t->left) + count(t->right);
86     }
87     return counttree;
88 }
89
90 int max(int a, int b){
91     if (a >= b){
92         return a;
93     }
94     return b;
95 }
96
97 int height(Tree t){
98     int heighttree = -1;
```



```

99     if (t != NULL) {
100         heighttree = 1 + max(height(t->left), height(t->right));
101     }
102     return heighttree;
103 }
104
105 int balance (Tree t){ // calculates the difference between left
and right
106     int diff = 0;
107
108     if (t != NULL) {
109         diff = count(t->left) - count(t->right);
110         if (diff < 0) {
111             diff = -diff;
112         }
113     }
114     return diff;
115 }
116
117 void freeTree(Tree t) { // free in postfix fashion
118     if (t != NULL) {
119         freeTree(t->left);
120         freeTree(t->right);
121         free(t);
122     }
123     return;
124 }

```

The output is:

```

      1
     / \
    2   3
   / \
  4   5
 / \
6   7
/
8

```

Balance = 5
Height = 6
Count = 8

Deleting a node from a BST

Deletion is harder than insertions. We could:

- find the node to be deleted
- unlink the node from its parent

But what do we do with the deleted node's children?

Easy option, don't delete, just mark the node as deleted

- Future searches ignore this item
- Problem? Tree can become full of 'deleted nodes'!

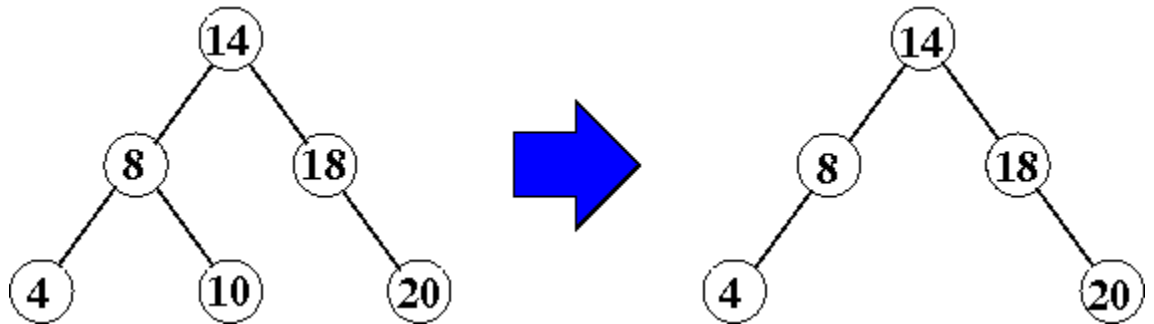
Hard option has 3 cases:

1. node is a **leaf**
 - there are no children, so unlink node from parent
2. node has **1 child**

- simply replace the node by its child
- 3. node has **2 children**
 - we need to rearrange the tree in some way

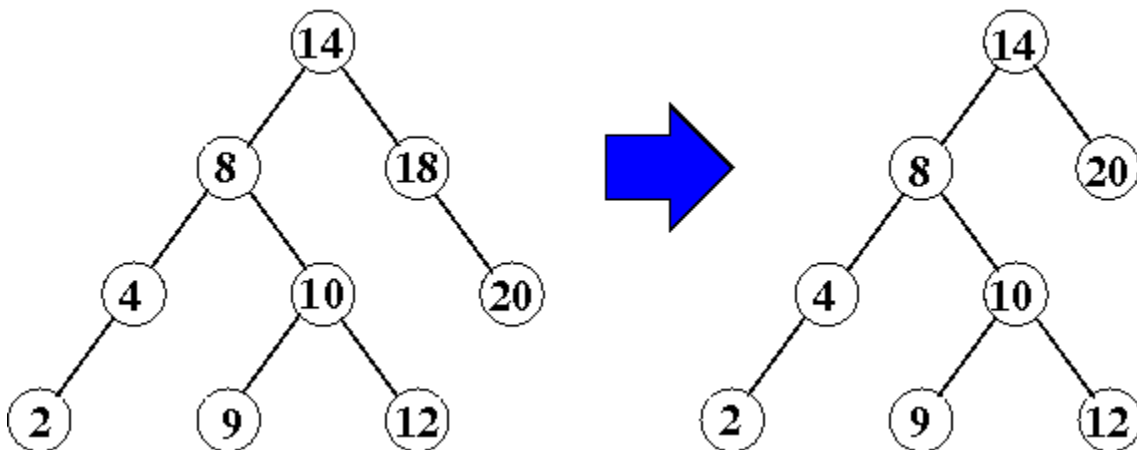
node is a leaf

- Delete 10 from this tree



node has 1 child

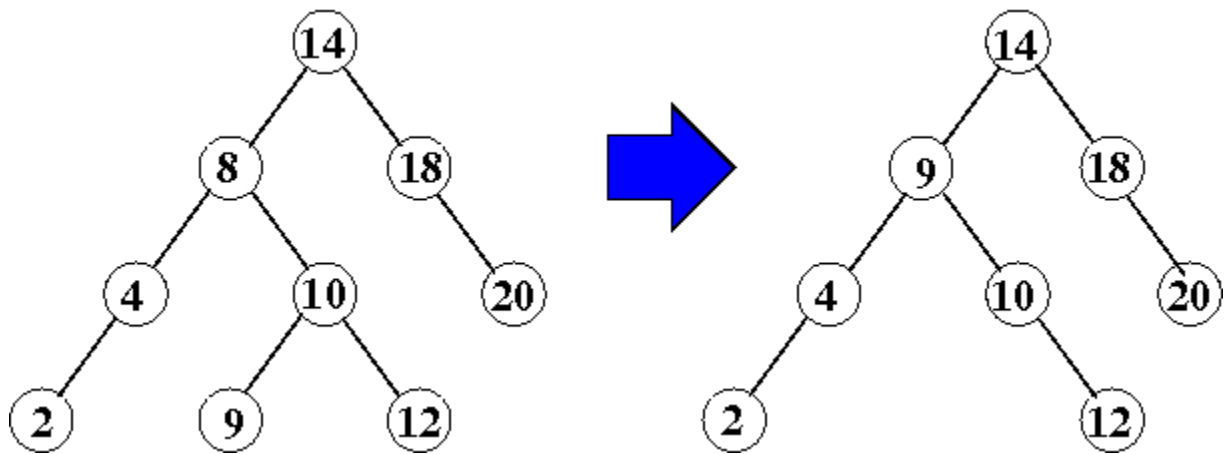
- Delete 18 from this tree



node has 2 children

- Delete 8 from this tree
- Join trees by replacing the node to be deleted with a 'descendant' of the deleted node
 - the descendant is either:
 1. the **right child's** Deepest Left-Most Descendent (DLMD)
 - this node has the 'smallest' larger value than the deleted node
 2. the **left child's** Deepest Right-Most Descendent (DRMD)
 - this node has the 'largest' smaller value than the deleted node

We use the first strategy here.



Notice:

1. the DLMD of the right child is 9 and we replaced 8 by this node
2. we could have chosen instead the DRMD of the left child is 4
 - ... which is 4, because 4 has no right descendents

Delete a node code

Here is the code that deletes an arbitrary node in a BST with 0, 1 or 2 children:

切换行号显示

```

1 Tree deleteTree(Tree t, int i){ // delete node with value 'v'
2     if (t != NULL) {
3         if (v < t->data) {
4             t->left = deleteTree(t->left, v);
5         }
6         else if (v > t->data) {
7             t->right = deleteTree(t->right, v);
8         }
9         else { // v == t->data, so the node 't' must be deleted
10            // next fragment of code violates style, just to make logic
clear
11            Tree n; //
temporary
12            if (t->left==NULL && t->right==NULL) n=NULL; // 0
children
13            else if (t->left ==NULL) n=t->right; // 1
child
14            else if (t->right==NULL) n=t->left; // 1
child
15            else n=joinDLMD(t->left,t-
>right);
16            free(t);
17            t = n;
18        }
19    }
20    return t;
21 }

```

切换行号显示

```

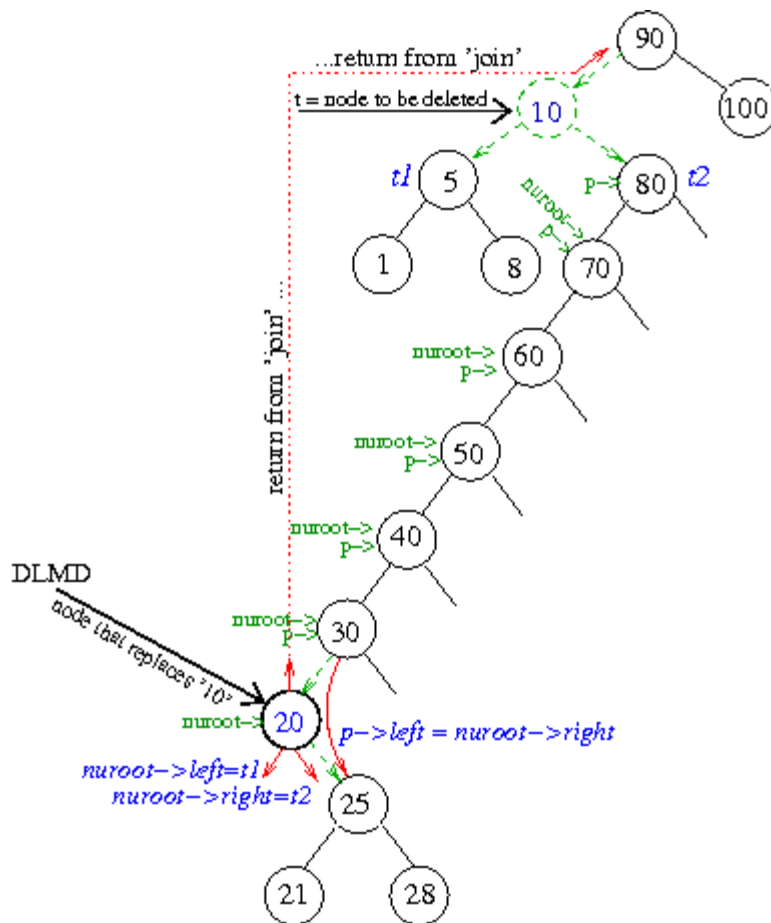
1 // Joins t1 and t2 with the deepest left-most descendent of t2 as
new root.
2 Tree joinDLMD(Tree t1, Tree t2){
3     Tree nurroot;
4     if (t1 == NULL) { // actually should never happen
5         nurroot = t2;
6     }

```

```

7   else if (t2 == NULL) {           // actually should never happen
8       nuroot = t1;
9   }
10  else {                           // find the DLMD of the right subtree
t2
11      Tree p = NULL;
12      nuroot = t2;
13      while (nuroot->left != NULL) {
14          p = nuroot;
15          nuroot = nuroot->left;
16      }                             // nuroot is the DLMD, p is its
parent
17      if (p != NULL) {
18          p->left = nuroot->right; // give nuroot's only child to
p
19          nuroot->right = t2;      // nuroot replaces deleted node
20      }
21      nuroot->left = t1;          // nuroot replaces deleted node
22  }
23  return nuroot;
24 }

```



BST operations complexity analysis

Cost of searching:

- Best case: key is at root: $O(1)$
- Worst case: key is not in BST: search the height of the tree
 - Balanced trees: $O(\lg n)$
 - Degenerate trees: $O(n)$
- Average case: key is in middle
 - Balanced tree: $O(\lg n)$

- Degenerate trees: $O(n)$

Insertion and deletion:

- Always traverse height of tree
 - Balanced tree: $O(\lg n)$
 - Degenerate tree: $O(n)$

Summary of operations

- A Binary Search Tree is an ordered binary tree:
 - *left child < parent < right child*
- Searching in BSTs
 - *int searchTree(Tree, int)*
- Inserting nodes into BSTs
 - *Tree createNode()*
 - *Tree insertTree(Tree, int)* recursive
 - *Tree insertTreeI(Tree, int)* iterative
 - *void freeTree()*
- Printing BSTs
 - laying a tree down 'on its side'
 - *printTree()*
- Characteristics of BSTs
 - *int height(Tree)*
 - *int count(Tree)*
 - *int balance(Tree)*
- Deletion from BSTs
 - *Tree deleteTree(Tree, int)* recursive
 - *Tree joinDLMD(Tree, Tree)* replace the node to be deleted with the DLMD
- Complexity analysis of BSTs

BinarySearchTrees (2019-08-05 11:53:06由AlbertNymeyer编辑)