

目录

1. Data structures: memory allocation
 1. sizeof structs
 1. Compile "padding.c" with "gcc"
 2. Compile "padding.c" with "dcc"
 3. Language definition says nothing about memory allocation

Data structures: memory allocation

A character is stored in a byte (8 bits)

- a *char* takes 1 byte (8 bits)
- an *int* takes 4 bytes (32 bits)
- ***how much memory does a pointer need?***

Commonly, computers are 32 bit

- means they support 32-bit addresses
 - an address is hence 4 bytes (same as an *int*)
- 4GB memory can be addressed

More modern computers are 64 bit

- addresses are hence 64 bits (= 8 bytes)
- *chars* are (still) 1 byte
- *ints* are (still) 4 bytes

Consider the program

切换行号显示

```
1 #include <stdio.h>
2
3 int main(void)
4 {
5     int i;
6     char c;
7
8     printf("i and c are uninitialised\n");
9     printf("&i = %p\n", (void*)&i);
10    printf("&c = %p\n", (void*)&c);
11
12    printf("sizeof(&i) = %ld, sizeof(i) = %ld\n", sizeof(&i),
sizeof(i));
13    printf("sizeof(&c) = %ld, sizeof(c) = %ld\n", sizeof(&c),
sizeof(c));
14
15    printf("now give i and c values\n");
16    i = 58;
17    c = 'X';
```

```

18  printf("i = %d\n", i);
19  printf("c = %c\n", c);
20  return 0;
21 }

```

Aside:

- why the cast (void *) in the program?

This program has two variables *i* and *c*.

- the address of both are taken and printed
- the *sizeof()* each variable, and their addresses are then printed
- the variables are then initialised
- their values are then printed

The output is

```

i and c are uninitialised
&i = 0x7ffd8383dc00
&c = 0x7ffd8383dc10
sizeof(&i) = 8, sizeof(i) = 4
sizeof(&c) = 8, sizeof(c) = 1
now give i and c values
i = 58
c = X

```

Notice:

- chars are 1 byte, ints are 4 bytes, and addresses are 8 bytes long
 - ... 8 bytes because it is a 64-bit machine (the address space)
 - we normally do not worry about the size of data

sizeof structs

Consider the following program:

切换行号显示

```

1  // padding.c
2  #include <stdio.h>
3  #include <stdlib.h>
4
5  #define SIZEBYTES 5
6  #define S(x)      sizeof(x) // this is a C macro that is
substituted by the preprocessor
7                                // I use it in the program just to
save line-length space
8  struct s1 {
9      char  c;
10     int   i;
11     char *s;
12 };
13 struct s2 {
14     char  c;
15     char *s;
16     int   i;

```

```

17 };
18
19 int main() {
20
21     struct s1 a;
22     struct s2 b;
23
24     printf("sizeof(a.c)=%lu, sizeof(a.i)=%lu, sizeof(a.s)=%lu\n",
S(a.c), S(a.i), S(a.s));
25     printf("sizeof(b.c)=%lu, sizeof(b.s)=%lu, sizeof(b.i)=%lu\n",
S(b.c), S(b.s), S(b.i));
26
27     printf("sizeof(a) = %lu, sizeof(b) = %lu\n\n", sizeof(a),
sizeof(b));
28
29     a.s = malloc(SIZEBYTES); // returned value not checked
30     printf("&a.c = %p\n", (void*)&a.c);
31     printf("&a.i = %p\n", (void*)&a.i);
32     printf("&a.s = %p\n\n", (void*)&a.s);
33
34     b.s = malloc(SIZEBYTES); // returned value not checked
35     printf("&b.c = %p\n", (void*)&b.c);
36     printf("&b.s = %p\n", (void*)&b.s);
37     printf("&b.i = %p\n", (void*)&b.i);
38
39     return 0;
40 }

```

Compile "padding.c" with "gcc"

The output will look like:

```

sizeof(a.c)=1, sizeof(a.i)=4, sizeof(a.s)=8
sizeof(b.c)=1, sizeof(b.s)=8, sizeof(b.i)=4
sizeof(a) = 16, sizeof(b) = 24

&a.c = 0x7fadedface30
&a.i = 0x7fadedface34
&a.s = 0x7fadedface38

&b.c = 0x7fadedface10
&b.s = 0x7fadedface18
&b.i = 0x7fadedface20

```

Notice:

- both structs contain 3 fields: a char, a pointer and an int
 - total is $1+8+4 = 13$ bytes
- order of fields is different in the two structs
- fields have addresses in declaration order (low to high)
- structs are in reverse declaration order (high to low)

But:

- var *a* uses 16 bytes
- var *b* uses 24 bytes
- *why different? why larger?*

Let's examine the fields in variable *a*:

- *a.c* (1 byte) occupies 30, so **bytes 31...33 are padding (3 bytes)**
- *a.i* (4 bytes) occupies 34...37
- *a.s* (8 bytes) occupies 38 ..3f
- Total: 13 bytes actual data + 3 bytes padding

Let's examine the fields in variable *b*:

- *b.c* (1 byte) occupies 10, so **bytes 11..17 are padding (7 bytes)**
- *b.s* (8 bytes) occupies 18..1f
- *b.i* (4 bytes) occupies 20..23, so **bytes 24..27 are padding (4 bytes)**
- Total: 13 bytes actual data + 11 bytes padding

Why is the compiler inserting padding?

It is a 64-bit address computer. A 'word' is hence 8 bytes (i.e. 64 bits)

- variable *a*
 - the *char* is padded to an *int* boundary (3 bytes inserted)
 - the *char* and the *int* share the same word
 - address *char* * naturally occupies one word
 - the *struct* has finished on a word boundary
- variable *b*
 - the *char* is padded to a word boundary (7 bytes inserted)
 - address *char* * naturally occupies one word
 - the *int* is padded to a word boundary (4 bytes inserted)
 - the *struct* has finished on a word boundary

Between *a* and *b* there is also padding

- *b* finishes at 0x7fadedface27, *a* starts at 0x7fadedface30: 8 bytes inserted

Note:

- none of the variables have been defined (no values assigned)
- changing *SIZEBYTES* will not change the output above
 - *Question: do you know why?*

Compile "padding.c" with "dcc"

The output is:

```
sizeof(a.c)=1, sizeof(a.i)=4, sizeof(a.s)=8
sizeof(b.c)=1, sizeof(b.s)=8, sizeof(b.i)=4
sizeof(a) = 16, sizeof(b) = 24

&a.c = 0x7fadedface60
&a.i = 0x7fadedface64
```

```
&a.s = 0x7fadedface68
&b.c = 0x7fadedface80
&b.s = 0x7fadedface88
&b.i = 0x7fadedface90
```

Note:

- appears the same as with *gcc* but ...
 - the memory order of *a* and *b* is in declaration order (not reversed)
 - *Question: what's the last memory address used by variable a?*
 - the padding between *a* and *b* is 16 bytes

Language definition says nothing about memory allocation

The C language definition does not define where variables are placed in memory

- memory allocation is dependent on compiler and underlying architecture (i.e. instruction set)
- different compilers generate different code for the same machine ('gcc' vs 'dcc' above)
- often a *space vs time* trade-off
 - every variable occupies 1 word optimises speed → poor space utilisation
 - 'packing' data in memory optimises space → poor performance (extra time cost to 'unpack' the data)

Lec01Sizeof (2019-06-14 11:02:32由AlbertNymeyer编辑)