

目录

1. Graph search
 1. Breadth-first versus Depth-first search
 2. Stack-Based Depth-First Search
 1. The helper ADT IOMem
 2. Testing Stack-Based Depth-First Search
 3. Performance
 3. Recursive Depth-First Search
 1. Testing recursive DFS
 2. Globally visited
 3. Testing a disconnected graph
 4. Breadth First Search
2. Applications of Depth-First Search
 1. Depth-First Search: Cycle Detection
 2. Eulerian cycles (using DFS)
 1. Eulerian Cycle Algorithm
 3. Path searching (using DFS)
 4. Reachability Analysis
3. Application of Breadth-First Search: path searching

Graph search

Searching a graph can have many aims:


- can I reach every vertex in the graph (is it connected)?
- is one vertex reachable starting from some other vertex?
- what is the shortest path from vertex v to w ?
- which vertices are reachable from a vertex? (transitive closure)
- is there a cycle that passes through all the graph? (*tour*)
- is there a tree that links all vertices? (*spanning tree*)
 - what is the *minimum* spanning tree?
- are two graphs "equivalent"? (*isomorphism*)

A search is almost never 'random': it uses an underlying strategy:


- depth-first search DFS
- breadth-first search BFS

Breadth-first versus Depth-first search

Example:



Depth-First Search



Breadth-First Search

Order is given by the 'red' labels

- in this example the label ordering is breadth-first (layer by layer)

DFS descends by selecting the first available unvisited node

- select 0
- adjacent {1,2,3}
 - select 1
 - adjacent {2, 4}
 - select 2
 - adjacent {3, 5}
 - select 3
 - adjacent {5, 6}
 - select 5
 - adjacent {4, 8}
 - select 4
 - adjacent {7}
 - select 7
 - adjacent {8}
 - select 8
 - adjacent {6}
 - select 6
 - adjacent {} no sites left unvisited

BFS descends by systematically visiting the nodes in order of level

- select 0
- adjacent {1,2,3}
 - select 1
 - adjacent {4}
 - select 2
 - adjacent {5}
 - select 3
 - adjacent {6}
 - select 4
 - adjacent {7}
 - select 5

- adjacent {8}
- select 6
- adjacent {}
 - select 7
 - adjacent {}
 - select 8
 - adjacent {} no sites left unvisited

These two 'strategies' actually use the same algorithm. They differ only in their use of data structure:

- DFS uses a stack
- BFS uses a queue

Here is the pseudo-algorithm for **Depth/Breadth**-first search:

```

push the root node onto a stack/queue
while (stack/queue is not empty) {
    pop a node from the stack/queue
    if (node is a goal node)
        return 'success'
    push all children of node onto the stack/queue
}
return 'failure'

```

If the aim is not to find a goal node, but to search the whole graph:

- leave out the conditional 'return' (i.e if (node is ...)
- return when complete

Stack-Based Depth-First Search


When searching we need to remember which nodes we've *visited*:

- to avoid cycles
- to make sure every node gets visited


Generally an array **visited[0 .. numVertices-1]** is used

- array indices correspond to vertices
- initialise all elements to -1, meaning unvisited
- when a vertex is visited, the index is set to its 'visit order' number
 - this is simply a 'count' that gets incremented each time a new node is visited

For example, here is the earlier graph again



Depth-First Search



Breadth-First Search

The *visited* array starts as $\{-1, -1, -1, -1, -1, -1, -1, -1, -1\}$

We select the root **0** first

adjacent	visit	resulting visited array
<i>any node</i>	0	$\{ \mathbf{0}, -1, -1, -1, -1, -1, -1, -1, -1 \}$
1 2 3	1	$\{ 0, \mathbf{1}, -1, -1, -1, -1, -1, -1, -1 \}$
0 2 4	2	$\{ 0, 1, \mathbf{2}, -1, -1, -1, -1, -1, -1 \}$
0 1 3 5	3	$\{ 0, 1, 2, \mathbf{3}, -1, -1, -1, -1, -1 \}$
0 2 5 6	5	$\{ 0, 1, 2, 3, -1, \mathbf{4}, -1, -1, -1 \}$
2 3 4 8	4	$\{ 0, 1, 2, 3, \mathbf{5}, 4, -1, -1, -1 \}$
1 5 7	7	$\{ 0, 1, 2, 3, 5, 4, -1, \mathbf{6}, -1 \}$
4 8	8	$\{ 0, 1, 2, 3, 5, 4, -1, 6, \mathbf{7} \}$
5 6 7	6	$\{ 0, 1, 2, 3, 5, 4, \mathbf{8}, 6, 7 \}$

Let's try a different starting vertex: this time start at vertex **5**:

adjacent	visit	resulting visited array
<i>any node</i>	5	$\{-1, -1, -1, -1, -1, \mathbf{0}, -1, -1, -1\}$
2 3 4 8	2	$\{-1, -1, \mathbf{1}, -1, -1, 0, -1, -1, -1\}$
0 1 3 5	0	$\{ \mathbf{2}, -1, 1, -1, -1, 0, -1, -1, -1 \}$
1 2 3	1	$\{ 2, \mathbf{3}, 1, -1, -1, 0, -1, -1, -1 \}$
0 2 4	4	$\{ 2, 3, 1, -1, \mathbf{4}, 0, -1, -1, -1 \}$
1 5 7	7	$\{ 2, 3, 1, -1, 4, 0, -1, \mathbf{5}, -1 \}$
4 8	8	$\{ 2, 3, 1, -1, 4, 0, -1, 5, \mathbf{6} \}$
5 6 7	6	$\{ 2, 3, 1, -1, 4, 0, \mathbf{7}, 5, 6 \}$
3 8	3	$\{ 2, 3, 1, \mathbf{8}, 4, 0, 7, 5, 6 \}$


The array *visited*[*i*] here is the **depth-first order**

- It says: { 2nd, 3rd, 1st, 8th, 4th, 0th, 7th, 5th, 6th}

The visited array indicates the order of the search.

Can anything go wrong during the traversal?

- *Yes, we can hit a deadend!*



choice	visit	resulting visited array
<i>any node</i>	0	{ 0 , -1, -1, -1, -1, }
1 2	1	{ 0, 1 , -1, -1, -1 }
0 3	3	{ 0, 1, -1, 2 , -1 }
1 4	4	{ 0, 1, -1, 2, 3 }
	<i>finished?</i>	

- 4 is a leaf node, we can go no further
- **there is still an unvisited vertex in the array**

How do we 'find' it?

- we need to **backtrack**
 - we go back to vertex 0, and then visit vertex 2
 - this is also a leaf node
 - ... but all nodes have been visited, so we are really finished this time

Final DFS path is visited = {0, 1, 4, 2, 3}

So we cannot expect DFS to visit every vertex in a single forward traversal

- we sometimes need to *backtrack*


But how do we backtrack?

- we use a stack!
 - vertices are pushed onto the stack when we have 1 or more adjacent vertices to visit
 - to actually visit a vertex, we simply pop it from the stack
- *only when the stack is empty have we visited everyone!*


Using a stack in DFS means:

- when we *visit*, we *pop* the next vertex off the stack
- after a visit, we *push* the adjacent vertices onto the stack
- when we land on a leaf node, we cannot *push* any nodes onto the stack
 - we then *pop* a vertex instead
 - ... this is backtracking (to an earlier vertex)
- only when the stack is empty have we visited every vertex

Consider the above graph again:



The following stack operations are carried out:



Code:

切换行号显示

```

1 // dfsStack.c: traverse a graph using DFS and stacking
  (graph may be disconnected)
2 // Compile using:
3 //      dcc -o dfsStack dfsStack.c IOmem.c GraphAM.c
Quack.c
4 //
5 #include <stdio.h>
6 #include <stdlib.h>
7 #include "Graph.h"
8 #include "Quack.h"
9 #include "IOmem.h"
10
11 #define STARTVERTEX 0    // start the depth-first search
  at this vertex
12
13 void dfsQuack(Graph, Vertex, int);
  
```

```

14
15 int main (void) {
16     int numV;
17     if ((numV = readNumV()) > 0) {
18         Graph g = newGraph(numV);
19         if (readBuildGraph(g)) {
20             showGraph(g);
21             dfsQuack(g, STARTVERTEX, numV);
22         }
23         g = freeGraph(g);
24         g = NULL;
25     }
26     else {
27         printf("Error in reading #number\n");
28         return EXIT_FAILURE;
29     }
30     return EXIT_SUCCESS;
31 }
32
33 void dfsQuack(Graph g, Vertex v, int numV) {
34     int *visited = mallocArray(numV);
35     Quack s = createQuack();
36     push(v, s);
37     showQuack(s);
38     int order = 0;
39     while (!isEmptyQuack(s)) {
40         v = pop(s);
41         if (visited[v] == UNVISITED) { // we visit only
unvisited vertices
42             printArray("Visited: ", visited, numV);
43             visited[v] = order++;
44             for (Vertex w = numV - 1; w >= 0; w--) { //
push adjacent vertices
45                 if (isEdge(newEdge(v,w), g)) { // ...
in reverse order
46                     push (w, s); // ...
onto the stack
47                 }
48             }
49         }
50         showQuack(s);
51     }
52     printArray("Visited: ", visited, numV);
53     free(visited);
54     return;
55 }

```

The helper ADT IOMem

Input/output and memory management is controlled by an ADT called **IOMem**. It's interface is:

切换行号显示

```
1 // IOmem.h
2 // Interface to IOmem ADT that reads input data, builds
and print graphs and manages memory.
3
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 int readNumV();           // read an int (numV) from
stdin
8 int readBuildGraph(Graph); // read int pairs from
stdin
9 int* mallocArray(int);    // malloc an array of
length int * sizeof(int)
10 void printArray(char *, int *, int); // print an int
array of length int
11
```

This ADT allows the amount of graph search code to be kept minimal.

We can now compile the graph search algorithm with the *Graph*, *Quack* and *IOmem* ADTs:

- we can either the *GraphAM* or *GraphAL* ADTs

```
gcc dfsStack.c IOmem.c GraphAM.c Quack.c
```

or

```
gcc dfsStack.c IOmem.c GraphAL.c Quack.c
```


- also a choice between the array-based ADT *Quack* and linked-list version *QuackLL*
- in total, 4 combinations of *Graph* and *Quack* ADTs possible!

Testing Stack-Based Depth-First Search

The input file we use is:

```
#5
0 1 0 2 1 3 3 4
```

which corresponds to the simple graph we saw before:



Executing *a.out* using this input file results in the following:

```

V=5, E=4
<0 1> <0 2>
<1 0> <1 3>
<2 0>
<3 1> <3 4>
<4 3>
Quack: <<0>>
Visited: {-1, -1, -1, -1, -1}
Quack: <<1, 2>>
Visited: {0, -1, -1, -1, -1}
Quack: <<0, 3, 2>>
Quack: <<3, 2>>
Visited: {0, 1, -1, -1, -1}
Quack: <<1, 4, 2>>
Quack: <<4, 2>>
Visited: {0, 1, -1, 2, -1}
Quack: <<3, 2>>
Quack: <<2>>
Visited: {0, 1, -1, 2, 3}
Quack: <<0>>
Quack: << >>
Visited: {0, 1, 4, 2, 3}

```

Here we see:

- the starting vertex 0 is pushed
- 0 is popped and its neighbours 1 and 2 are pushed
 - visited[0] = 0
- 1 is popped and its neighbours 0 and 3 are pushed
 - visited[1] = 1
- 0 is popped and ignored as it is in array *visited*
- 3 is popped and its neighbours 1 and 4 are pushed
 - visited[3] = 2
- 1 is popped and is ignored
- 4 is popped and its neighbour 3 is pushed
 - visited[4] = 3

- 7/22/2019, 7:34 PM

```

Quack: <<3, 4, 5, 6, 7, 5, 7>>
Quack: <<4, 5, 6, 7, 5, 7>>
Quack: <<5, 6, 7, 5, 7>>
Quack: <<6, 7, 5, 7>>
Quack: <<7, 5, 7>>
Visited: {0, -1, 1, 4, 3, 5, 2, -1}
Quack: <<0, 1, 4, 5, 7>>
Quack: <<1, 4, 5, 7>>
Visited: {0, -1, 1, 4, 3, 5, 2, 6}
Quack: <<7, 4, 5, 7>>
Quack: <<4, 5, 7>>
Quack: <<5, 7>>
Quack: <<7>>
Quack: << >>
Visited: {0, 7, 1, 4, 3, 5, 2, 6}

```

Performance

- number of pushes and pops
 - should be the same (stack is empty at the end)
- number of pushes of a vertex v = vertex degree of v
- total number of pushes
 - = sum of all the vertex degrees of vertices v in the graph

The sum of vertex degrees is equal to twice the number of edges.

- this means the complexity is linear in the number of edges, $O(E)$
 - **what does this mean? ...**
 - *how many edges are there?*
 - the worst case is a dense graph: $E = V*(V-1)/2$
 - so the complexity is quadratic in V : i.e. $O(V^2)$
 - if it is sparse, then it will be less than quadratic
- often said that DFS is *linear in the size of the graph* ...
 - ... where 'size' is the number of edges
 - ... which is another way of saying *quadratic in the number of vertices*

Recursive Depth-First Search

DFS above used our own stack to 'remember' which path it was traversing and do backtracking.

The system also has a stack, called a *call stack*, which is used to execute functions

- a function call causes a *function frame* to be pushed onto the call stack
 - ... upon a function return, the *frame* is popped off the call stack

This works even for recursive functions of course.

The call stack can be used instead of the 'stack' ADT we used above

- so when we compile we do not need the ADT quack

It works by recursion:

- a function *dfsR()* calls itself recursively ...
- in essence adjacent vertices are being *pushed* onto the system 'call' stack
 - the *for-loop* in *dfsR()* is over all all unvisited adjacent vertices
 - in the *for-loop*, *dfsR()* is called for every vertex
 - these calls *stack up* as you descend down the tree

切换行号显示

```

1 // dfsRec.c: traverse a graph using DFS (graph may be
disconnected)
2 // Compile using:
3 //     dcc -o dfsRec dfsRec.c IOMem.c GraphAM.c
4 //
5 #include <stdio.h>
6 #include <stdlib.h>
7 #include "Graph.h"
8 #include "IOMem.h"
9
10 void dfs(Graph, Vertex, int);
11 void dfsR(Graph, Vertex, int, int *, int *);
12
13 int main(void) {
14     int numV;
15     if ((numV = readNumV()) >= 0) {
16         Graph g = newGraph(numV);
17         if (readBuildGraph(g)) {
18             showGraph(g);
19             dfs(g, 0, numV); // DEPTH-FIRST SEARCH FROM
NODE 0
20         }
21         g = freeGraph(g);
22         g = NULL;
23     }
24     else {
25         return EXIT_FAILURE;
26     }
27     return EXIT_SUCCESS;
28 }
29
30 void dfs(Graph g, Vertex v, int numV) { // a 'wrapper'
for recursive dfs
31     int *visited = mallocArray(numV); // ... handles
disconnected graphs
32     int order = 0;
33     Vertex newv = v; // this is the
starting vertex
34     int allVis = 0; // assume not all
visited

```

```

35     while (!allVis) {                                // as long as
there are vertices
36         dfsR(g, newv, numV, &order, visited);
37         allVis = 1;                                // are all visited
now?
38         for (Vertex w = 0; w < numV && allVis; w++) { //
look for more
39             if (visited[w] == UNVISITED) {
40                 printf("Graph is disconnected\n"); // debug
41                 allVis = 0;                        // found an
unvisited vertex
42                 newv = w;                          // next loop dfsR
this vertex
43             }
44         }
45     }
46     printArray("Visited: ", visited, numV);
47     free(visited);
48     return;
49 }
50
51 void dfsR(Graph g, Vertex v, int numV, int *order, int
*visited) {
52     visited[v] = *order;                            // records the
order of visit
53     printf("Visiting vertex %d in order %d\n", v,
*order);
54     *order = *order+1;
55     for (Vertex w = 0; w < numV; w++) {
56         if (isEdge(newEdge(v,w), g) &&
visited[w]==UNVISITED) {
57             dfsR(g, w, numV, order, visited);
58         }
59     }
60     return;
61 }

```

Here the function *dfs()* is called by *main()*


- this function is a *wrapper*
- it does 'housekeeping' (initialising *order* and the array *visited*)
- it calls the recursive function *dfsR()*

Remember: in the stack version the main function called *dfsQuack()*

- it does 'housekeeping' (initialising *order* and the array *visited*)
- *pops* and *pushes* off/on the quack until the quack is empty

Testing recursive DFS

Let's run this recursive DFS on the graph we had above:



Remember, it is represented by the input data:

```
#8
0 2 0 5 0 7 2 6 1 7 4 7 4 6 4 3 3 5 4 5
```

Compiling:

```
dcc -o dfsRec dfsRec.c IOfem.c GraphAM.c
```

notice, no *Quack* ADT, and executing

```
V=8, E=10
<0 2> <0 5> <0 7>
<1 7>
<2 0> <2 6>
<3 4> <3 5>
<4 3> <4 5> <4 6> <4 7>
<5 0> <5 3> <5 4>
<6 2> <6 4>
<7 0> <7 1> <7 4>
Visiting vertex 0 in order 0
Visiting vertex 2 in order 1
Visiting vertex 6 in order 2
Visiting vertex 4 in order 3
Visiting vertex 3 in order 4
Visiting vertex 5 in order 5
Visiting vertex 7 in order 6
Visiting vertex 1 in order 7
Visited: {0, 7, 1, 4, 3, 5, 2, 6}
```

Comparing that with the stack version, the last lines shown below:

```
.
.
.
Visited: {0, -1, 1, 4, 3, 5, 2, 6}
Quack: <<7, 4, 5, 7>>
Quack: <<4, 5, 7>>
Quack: <<5, 7>>
Quack: <<7>>
Quack: << >>
```

```
Visited: {0, 7, 1, 4, 3, 5, 2, 6}
```

In summary:

- we've seen 2 versions of depth-first search:
 - an explicit stack version *dfsStack.c* that uses a *Quack* ADT
 - a call-stack version *dfsRec.c* that uses recursion

You could argue that the stack version:

- requires much less system resources (no recursion)
- does backtracking in an *iterative* manner, so will be much faster

Globally visited

Crucial in both versions is the array *visited[]* and integer variable *order*

- *visited[]* records unvisited vertices and the *order* of visiting
- stop cycles occurring (remember, we are dealing with graphs)

Almost always implemented as global variables.

- For example, in recursive DFS, *dfsRec.c*
 - *visited[]* and *order* are initialised in the 'wrapper'
 - if global variables are used, they

GraphSearchDFSglobal


Testing a disconnected graph

We saw in *dfsRec.c* that the recursion handles disconnected graphs

Let's check that.

```
#6
0 1 0 2 1 2 3 4 3 5 4 5
```

corresponding to:



and assuming the starting vertex is 0, then *dfsRec* produces:

```
V=6, E=6
<0 1> <0 2>
<1 0> <1 2>
<2 0> <2 1>
```

```

<3 4> <3 5>
<4 3> <4 5>
<5 3> <5 4>
Visiting vertex 0 in order 0
Visiting vertex 1 in order 1
Visiting vertex 2 in order 2
Graph is disconnected
Visiting vertex 3 in order 3
Visiting vertex 4 in order 4
Visiting vertex 5 in order 5
Visited: {0, 1, 2, 3, 4, 5}

```

Notice:


- the first (sub) graph's DFS search begins at vertex 0
- the second (sub) graph's DFS search begins at vertex 3

Breadth First Search

All adjacent vertices are visited before moving to another vertex

- each level of vertices is visited before the next level's vertices are considered

For example:



1. visit vertex 0
2. visit vertex 1 and 2
3. visit vertex 3
4. visit vertex 4

In essence, the vertices are processed ***in order*** (top to bottom, left to right)

- DFS used a stack:
 - we pushed all the adjacent vertices of a vertex onto a stack
 - so we would remember the vertices we need to 'still visit'
- BFS instead uses a queue:
 - we push all the adjacent vertices of a vertex onto a queue (not a stack)
 - but we visit them in the order they occur (as we must in a queue)
- the change in the *quack* version is trivial:

- just 4 changes

切换行号显示

```


1 void bfsQuack(Graph g, Vertex v, int numV) { //name
change
2     int *visited = mallocArray(numV);
3     Quack q = createQuack();
4     qush(v, q); //qush, not
push
5     showQuack(q);
6     int order = 0;
7     while (!isEmptyQuack(q)) {
8         v = pop(q);
9         if (visited[v] == UNVISITED) {
10             printf("\t\t\t ... visit %d\n", v);
11             visited[v] = order++;
12             Vertex w;
13             for (w = 0; w < numV; w++) { //vertex order
14                 if (isEdge(newEdge(v,w), g)) {
15                     qush(w, q); //qush, not
push
16                 }
17             }
18         }
19         showQuack(q);
20     }
21     printArray("Visited: ", visited, numV);
22     free(visited);
23     makeEmptyQuack(q);
24     return;
25 }

```

What about the graph we considered above for DFS represented by the input data

#8

0 2 0 5 0 7 2 6 1 7 4 7 4 6 4 3 3 5 4 5



Starting at vertex 0, what did DFS do:

- $Visited[] = \{0, 7, 1, 4, 3, 5, 2, 6\}$
- corresponds to the vertices: 0 2 6 4 3 5 7 1

What does *bfs()* do:

```

Quack: <<0>>                                <== start node
Visited: {-1, -1, -1, -1, -1, -1, -1, -1}
Quack: <<2, 5, 7>>                            <== 2,5,7 pushed
Visited: {0, -1, -1, -1, -1, -1, -1, -1}
Quack: <<5, 7, 0, 6>>                        <== 0,6 quashed
Visited: {0, -1, 1, -1, -1, -1, -1, -1}
Quack: <<7, 0, 6, 0, 3, 4>>                  <== 0,3,4 pushed
Visited: {0, -1, 1, -1, -1, 2, -1, -1}
Quack: <<0, 6, 0, 3, 4, 0, 1, 4>>            <== 0,1,4 pushed
Quack: <<6, 0, 3, 4, 0, 1, 4>>
Visited: {0, -1, 1, -1, -1, 2, -1, 3}
Quack: <<0, 3, 4, 0, 1, 4, 2, 4>>            <== etc
Quack: <<3, 4, 0, 1, 4, 2, 4>>
Visited: {0, -1, 1, -1, -1, 2, 4, 3}
Quack: <<4, 0, 1, 4, 2, 4, 4, 5>>
Visited: {0, -1, 1, 5, -1, 2, 4, 3}
Quack: <<0, 1, 4, 2, 4, 4, 5, 3, 5, 6, 7>>
Quack: <<1, 4, 2, 4, 4, 5, 3, 5, 6, 7>>
Visited: {0, -1, 1, 5, 6, 2, 4, 3}
Quack: <<4, 2, 4, 4, 5, 3, 5, 6, 7, 7>>
Quack: <<2, 4, 4, 5, 3, 5, 6, 7, 7>>
Quack: <<4, 4, 5, 3, 5, 6, 7, 7>>
Quack: <<4, 5, 3, 5, 6, 7, 7>>
Quack: <<5, 3, 5, 6, 7, 7>>
Quack: <<3, 5, 6, 7, 7>>
Quack: <<5, 6, 7, 7>>
Quack: <<6, 7, 7>>
Quack: <<7, 7>>
Quack: <<7>>
Quack: <<>>
Visited: {0, 7, 1, 5, 6, 2, 4, 3}

```

In the first few lines of this output we see:

- vertex 0 is quashed, then popped
- vertices 2, 5, 7 are quashed, then successively
 - 2 is popped
 - 5 is popped
 - 7 is popped
- 'quashing' means everything gets added to the end of the quack
- the order that vertices are visited is:
 - 0 2 5 7 6 3 4 1
 - note:
 - 0 is a level-0 vertex
 - 2 5 7 are level-1
 - 6 3 4 1 are level-2
 - no vertex is more than a path length of 2 away from the starting vertex

Applications of Depth-First Search

Depth-First Search: Cycle Detection

A graph that does not contain a cycle is called a tree, so

- asking whether a graph contains no cycles is equivalent to
- asking whether a graph is a tree

The *recursive* DFS algorithm, *dfsR()* we saw earlier is:

切换行号显示

```

1 void dfsR(Graph g, Vertex v, int numV, int *order, int
*vvisited) {
2     visited[v] = *order;
3     *order = *order+1;
4     Vertex w;
5     for (w=0; w < numV; w++) {
6         if (isEdge(g, newE(v,w)) && visited[w]==UNVISITED)
{
7             dfsR(g, w, numV, order, visited);
8         }
9     }
10    return;
11 }
```

To search for a cycle, the function *main()* calls *searchForCycle()*:

- which does housekeeping and in turn
- calls the recursive function *hasCycle()*

The function *hasCycle()* is a simple modification of *dfsR()*. The changes are:

- introduce a variable *found* to terminate the search if a cycle is found
 - this is a sort of early-exit
 - this variable must be passed up the recursive calls
- separate the 2 conditions in the for-loop
 - if *w* is adjacent to *v* then
 - if *w* is UNVISITED then recurse (in other words, keep searching)
 - else *w* has been visited before and **we have a cycle**
 - (assuming we are not going backwards along the same path)
 - to avoid going backwards, we pass 2 vertices to *hasCycle()*

The program is as follows:

切换行号显示

```



1 void searchForCycle(Graph g, int v, int numV) {
2     int *mallocArray(int numV) {
3         // as before
4     }
```

```

5     void showArray(int *vis, int numV) {
6         // as before
7     }
8     int *visited = mallocArray(numV);
9     int order = 0;
10
11     if (hasCycle(g, numV, v, v, &order, visited)) {
12         printf("found a cycle\n");
13     }
14     else {
15         printf("no cycle found\n");
16     }
17     showArray(visited, numV);
18     free(visited);
19     return;
20 }
21
22 int hasCycle(Graph g, int numV, Vertex fromv, Vertex v,
int *order, int *visited) {
23     int retval = 0;
24     visited[v] = *order;
25     *order = *order+1;
26     Vertex w;
27     for (w=0; w<numV && !retval; w++) {
28         if (isEdge(g, newE(v,w))) {
29             if (visited[w]==UNVISITED) {
30                 printf("traverse edge %d-%d\n", v, w);
31                 retval = hasCycle(g, numV, v, w, order,
visited);
32             }
33             else {
34                 if (w != fromv) { // exclude the vertex
we've just come from
35                     printf("traverse edge %d-%d\n", v, w);
36                     retval = 1;
37                 }
38             }
39         }
40     }
41     return retval;
42 }

```

If we input the graph on the left below:



a program that calls this function will output:

```
found a cycle
```

Alternatively, if the input is the graph on the right, it will output:

```
no cycle found
```

Eulerian cycles (using DFS)

An Eulerian **path** in a graph is a path that includes every edge exactly once

- this path may include many visits to the same vertex

In turn, an Eulerian **cycle** is a special case of a Eulerian path in which the start and end points are the same

Animation of the *Konigsburg Bridge Problem*

- <http://www.youtube.com/watch?v=3bUvjajakGM>

Graph animation of finding an Euler cycle

- <http://www.cs.sunysb.edu/~skiena/combinatorica/animations/euler.html>

Well-known properties of Eulerian paths and cycles:

```
A connected graph has an Eulerian cycle if all its vertices  
have even degree.
```

```
A connected graph has no Eulerian cycle if any of its vertices  
has odd degree.
```

```
A connected graph has an Eulerian path if it has exactly 2  
vertices of odd degree, and all the rest have even degree.
```

A connected graph has no Eulerian path if it has more than 2 vertices of odd degree.

So, a graph with more than 2 vertices of odd degree has no Eulerian cycle or path.


If a graph contains a Eulerian path, it cannot contain a Eulerian cycle, and vice versa.

A graph that consists of vertices all with even degree is often called an `''''Eulerian graph''''`


There is a simple algorithm to find an Eulerian cycle in an Eulerian graph

- it uses a DFS, hence requires a stack (to backtrack when you reach a deadend).


Examples of Eulerian graphs:




`'box'`



`'bow'`



`'propbows'`



`'concsquares'`

Eulerian Cycle Algorithm

Assume graph is connected and is Eulerian.

1. Read the Eulerian graph and initialise a stack.

2. Choose any vertex v and push v .
3. While the stack is not empty:
 - if the top of stack vertex has one or more adjacent vertices
 - select arbitrarily the largest vertex w
 - push w
 - remove edge $v-w$
 - else pop the top vertex and print it

The sequence of vertices that is printed is an Eulerian cycle. How does the algorithm work?

- we traverse the graph by pushing vertices on the stack and removing edges that lead to them
 - the vertex that is pushed is the neighbour of the previous vertex that is pushed
- we continue to push vertices as long as we can
- if we have a vertex on the stack with no neighbours, we start to pop and print

The function `findEulerCycle()` below implements the algorithm:

切换行号显示

```

1 void findEulerCycle(Graph g, int numV, Vertex startv) {
2     Quack s = createQuack();
3     printf("Eulerian cycle: ");
4
5     push(startv, s);
6     while (!isEmptyQuack(s)) {
7         Vertex v = pop(s); // v is the top of stack vertex
and ...
8         push(v, s);        // ... the stack has not
changed
9         Vertex w;
10        if ((w = getAdjacent(g, numV, v)) >= 0) {
11            push(w, s);    // push a neighbour of v onto
stack
12            removeE(g, newE(v, w)); // remove edge to
neighbour
13        }
14        else {
15            w = pop(s);
16            printf("%d ", w);
17        }
18    }
19    putchar('\n');
20 }
21
22 Vertex getAdjacent(Graph g, int numV, Vertex v) {
23     // returns the Largest Adjacent Vertex if it exists,
else -1
24     Vertex w;
25     Vertex lav = -1; // the adjacent vertex
26     for (w=numV-1; w>=0 && lav== -1; w--) {
27         Edge e = newE(v, w);

```

```

28         if (isEdge(g, e)) {
29             lav = w;
30         }
31     }
32     return lav;
33 }

```

For example: for the **box** graph above:

```

push 0
push 3 and remove 0-3
push 2 and remove 3-2
push 1 and remove 2-1
push 0 and remove 1-0 <-- at this point all the edges have
been removed, and 0 is isolated
pop 0 and print 0
pop 1 and print 1
pop 2 and print 2
pop 3 and print 3
pop 0 and print 0

```

- **Eulerian cycle: 0 1 2 3 0**
- *It is not obvious why you need a stack here: 5 pushes were followed by 5 pops.*

We can reach a deadend during the traversal

- if the top vertex on the stack has no adjacent vertices, the traversal has reached a deadend
 - (remember that we are removing edges as we traverse the graph)
- but there may be vertices on the stack that **do** have branches to vertex neighbours
 - if there is one, then a vertex neighbour is pushed and the traversal continues
 - this is a process of **back-tracking** of course
- the process stops when branches have been taken
 - this will happen when all edges have been removed
 - every vertex on the stack will have been popped and printed
 - we may see the same vertex many times, but each edge is traversed just once

In the next example, for the **bow** graph above, we see backtracking in action:

```

push 0
push 2 and remove 0-2
push 6 and remove 2-6
push 1 and remove 6-1
push 0 and remove 1-0
pop 0 and print 0
pop 1 and print 1
pop 6 and print 6
push 5 and remove 2-5
push 4 and remove 5-4
push 3 and remove 4-3
push 2 and remove 3-2

```



```

pop 2 and print 2
pop 3 and print 3
pop 4 and print 4
pop 5 and print 5
pop 2 and print 2
pop 0 and print 0

```

- **Eulerian cycle: 0 1 6 2 3 4 5 2 0**

You may need to backtrack many times of course. Consider the **propbows** graph above:

```

push 0
pushed 11 and remove edge 0-11
pushed 10 and remove edge 11-10
pushed 9 and remove edge 10-9
pushed 6 and remove edge 9-6
pushed 5 and remove edge 6-5
pushed 9 and remove edge 5-9
pushed 0 and remove edge 9-0
pushed 8 and remove edge 0-8
pushed 7 and remove edge 8-7
pushed 0 and remove edge 7-0
popped 0 and print 0
popped 7 and print 7
popped 8 and print 8
popped 0 and print 0
popped 9 and print 9
popped 5 and print 5
popped 6 and print 6
popped 9 and print 9
pushed 4 and remove edge 10-4
pushed 3 and remove edge 4-3
pushed 10 and remove edge 3-10
popped 10 and print 10
popped 3 and print 3
popped 4 and print 4
popped 10 and print 10
pushed 2 and remove edge 11-2
pushed 1 and remove edge 2-1
pushed 11 and remove edge 1-11
popped 11 and print 11
popped 1 and print 1
popped 2 and print 2
popped 11 and print 11
popped 0 and print 0

```

- **Eulerian cycle: 0 7 8 0 9 5 6 9 10 3 4 10 11 1 2 11 0**

Path searching (using DFS)

You could want to search for a path between two given vertices:

- the starting vertex we already have
- add the goal vertex as a parameter and test 'is there a path' in the *dfs()* function.

切换行号显示

```

1  int isPath(Graph g, Vertex v, Vertex goalv, int numV,
int *order, int *visited) {
2      int found = 0;
3      visited[v] = *order;
4      *order = *order+1;
5      if (v == goalv) {
6          found = 1;
7      }
8      else {
9          Vertex w;
10         for (w=0; w < numV && !found; w++) {
11             if (isEdge(g, newE(v,w))) {
12                 if (visited[w] == UNVISITED) {
13                     found = isPath(g, w, goalv, numV, order,
visited);
14                     printf("path %d-%d\n", w, v);
15                 }
16             }
17         }
18     }
19     return found;
20 }
```

- this function does a DFS traversal, exactly like *dfsR()* ...
 - ... but at the same time, it searches for a path to a vertex *goalv*

So, instead of the call to *dfsR()* (in *dfs()*):

切换行号显示

```

1      dfsR(g, v, numV, &order, visited);
```

we call the function *isPath()*:

切换行号显示


```

1      #define STARTV 0
2      #define GOALV 3
3      if (isPath(g, STARTV, GOALV, numV, &order, visited))
{ //notice the STARTV and GOALV arguments
4          printf("found path\n");
5      }
6      else {
7          printf("no path\n");
8      }
```

- Here a *#define* is used to name the start and goal vertices: this could better be input

interactively of course.

If we input the graph:



then the output is:



```
path 3-4
path 4-6
path 6-2
path 2-0
found path
Visited: {0, -1, 1, 4, 3, -1, 2, -1}
```

- Here we can see which vertices were visited

Reachability Analysis

In many problems we are interested in knowing which vertices are *reachable* from some start vertex.

- for example, in the graphs:



some of the vertices are unreachable from the start vertex 0, others are not

- if the graph is undirected, it is obviously disconnected
- (if the graph is directed then it may not be disconnected of course)

A DFS algorithm can be used to find all the reachable vertices

- simply run the algorithm from the start vertex
- on conclusion, check the visited array
 - any vertex (except the start vertex) that is unvisited is unreachable

An alternative method is to use so-called 'fixed-point' computation.

1. initialise:
 - a reachable set comprising of just the start vertex
 - every other vertex is considered unreachable
2. check every unreachable vertex v
 - if there is an edge from a vertex in the reachable set to v
 - then add v to the reachable set
3. repeat the previous step until the reachable set does not change
 - if the reachable set does not change, terminate

When the set does not change, we have reached a 'fixed point'

- the set of vertices in the reachable set can be reached from the start vertex
- all other vertices cannot be reached

Example: consider the graph *squareline* above

- let R be the set of reachable vertices, and the start vertex be 0
 - initially $R = \{0\}$
- consider vertices 1..5
 - 1 is adjacent to 0, add to R
 - 2 is adjacent to 0, add to R
 - $R = \{0, 1, 2\}$
 - R has changed, so repeat
- consider vertices 3..5
 - 3 is adjacent to 1, add to R
 - $R = \{0, 1, 2, 3\}$
 - R has changed, so repeat
- consider vertices 4 and 5
 - neither vertex is adjacent to a vertex in R
 - R does not change
- terminate the algorithm

Notice that you do not need to use a stack here, or recursion. There is no backtracking.

Application of Breadth-First Search: path searching

You could want to search for a path between two given vertices, *start* and *goal* say:

- the starting vertex we already have
- add the goal vertex as a parameter and test for it in the *bfs()* function

But, we don't follow any 'path' during BFS (as we did in DFS): we traverse by level:

- whenever we 'visit' a node, we must remember its parent
 - we store the parent of each vertex in an array called *parent[]*

- we hence need 2 arrays, *visited[]* and *parent[]*
- when we find the goal node, we're finished
 - the path 'backwards' to the goal node will be stored in *parent[]*
- we print the path from the *start* to the *goal* stored in *parent[]*

切换行号显示

```

1 void searchPath(Graph g, Vertex start, Vertex goal, int
numV) {
2     int *visited = mallocArray(numV);
3     int *parent = mallocArray(numV);
4     Quack q = createQuack();
5     qush(start, q);
6     int order = 0;
7     visited[start] = order++;
8     int found = 0;
9     while (!isEmptyQuack(q) && !found) {
10         Vertex x = pop(q);
11         for (Vertex y = 0; y < numV && !found; y++) {
12             if (isEdge(newEdge(x,y), g)) { // for
adjacent vertex y ...
13                 if (visited[y] == UNVISITED) { // ... if y
is unvisited ...
14                     qush(y, q); // ...
queue y
15                     printf("\t\t\t ... trying edge %d-%d\n",
x, y);
16                     visited[y] = order++; // y is now
visited
17                     parent[y] = x; // y's parent
is x
18                     if (y == goal) { // if y is the
goal ...
19                         found = 1; // ...
SUCCESS! now get out
20                     }
21                 }
22             }
23         }
24     }
25     if (found) {
26         printf("SHORTEST path from %d to %d is ", start,
goal);
27         printPath(parent, numV, goal); // an extern
function
28         putchar('\n');
29     }
30     else {
31         printf("no path found\n");
32     }
33     printArray("Visited: ", visited, numV);
34     printArray("Parent : ", parent, numV);


```

```

35     free(visited);
36     free(parent);
37     makeEmptyQuack(q);
38     return;
39 }

```

If we input the graph:



then the output is:

```

SHORTEST path from 0 to 6 is 6<--2<--0

```

The array *parent[]* stores the parent of every visited node in the graph, but:

- the start vertex has no parent
- unvisited nodes have no parents

To print the path:

- we print *goal*
- we print *parent[goal]*
- we print *parent[parent[goal]]*
- we print *parent[parent[parent[goal]]]*
- we print *parent[parent[parent[parent[goal]]]*
- ...
- until we find the *start* vertex

That's what the function below does:

切换行号显示

```

1 void printPath(int parent[], int numV, Vertex v) {
2     printf("%d", v);
3     if (0<=v && v<numV) {
4         Vertex p = parent[v];
5         while (p != UNVISITED) {
6             printf("<--%d", p);
7             p = parent[p];
8         }
9     }
10    else {

```

```
11         fprintf(stderr, "printPath: illegal vertex in\n");
12     }
13 }
```

As we start with *goal* and work backwards, we print the path in reverse direction.

GraphSearch (2019-07-22 17:59:26由AlbertNymeyer编辑)